To my beloved
Benedetta

Introduction

English version

Machine learning literature is exploding in size and complexity, but most solutions found are ad hoc, there is little communication between different subfields, and there is a large research debt. Category theory can solve these problems. [SGW21].

Talk about the origins of category theory and its "rise to power" as a common language that aims to unite different fields of knowledge.

Discuss the purpose of this work: a beginner-friendly survey of categorical approaches to neural networks, causal models, and interpretability.

Italian version

Traduzione italiana dell'introduzione.

Contents

Chapter 1

Parametric Optics for Gradient-Based Learning

General purpose deep learning libraries like PyTorch or TensorFlow are rich toolkits for gradient-based learning that allow machine learning practitioners to delegate low-level computations and automatic differentiation (AD) so that their focus can be directed to the high-level architectural design. Despite the unquestionable success of these libraries, their internal structure is not informed by a principled mathematical framework and thus, as mentioned in the introduction, it is highly dependent on *ad hoc* optimizations and empirical euristics (ADDREF). Consequently, adding new tools, composing existing ones, and parallelizing computations is not a straight forward process and often lead to unexpected results (ADDREF).

Hence, it would be auspicable to develop a mathematically structure compositional framework for gradient-based learning able to act as a bridge between low-level automatic differentiation and high level architectural specifications (ADDREF). Among the various attempts (ADDREF, ADDREF, ADDREF), a promising combination of differential categories, parametrization and optics has been recently proposed by [CCG+19], [CGG+22], [Gav24] and more as full-featured gradient-based framework able to challenge established tools. In this chapter, we illustrate such framwork and part of its mathematical foundations.

1.1 Categorical toolkit

Learning neural networks have two main properties: they depend on parameters and information flows through them bidirectionally (forward propagation and back propagation): any aspiring categorical model of gradient-based learning must take these two aspects into consideration. A number of authors (see e.g. [Gav24], and [CGG+22]) have proposed the **Para** construction as a categorical model of parameter dependence and various categories of optics as the right categorical abstraction for bidirectionality.

1

1.1.1 Actegories

Before we can deal with parametric maps, we need to find a way to glue input/output spaces to parameter spaces, so that such maps have well-defined domains. One common strategy is to provide the category at hand with a monoidal structure. However, monoidal products can only combine elements within the same underlying category. Since (co)parameters are often taken from spaces that are different in nature from the input and output spaces, a more general mathematical tool is needed: namely, actegories (see the survey [CG22] for a thorough treatment of the subject). Actegories are actions of symmetric monoidal categories on other categories. For brevity's sake, we will only give an incomplete definition (see [CG22] or [Gav24] for further information).

Definition 1 (Actegory). *Let* $(\mathcal{M}, I, \otimes)$ *be a strict symmetric monoidal category. A* $\mathcal{M}$*-actegory is a tuple* $(\mathcal{C}, \bullet, \eta, \mu)$*, where* $\mathcal{C}$ *is a category,* $\bullet : \mathcal{M} \times \mathcal{C} \to \mathcal{C}$ *is a functor, and* η *and* μ *are natural isomorphisms enforcing* $I \bullet C \stackrel{\eta_C}{\cong}$ *and* $(M \bullet (N \bullet C)) \stackrel{\mu_{M,N,C}}{\cong} (M \otimes N) \bullet C$*. The isomorphisms* η *and* μ *must also satisfy coherence conditions. If* η *and* μ *are identical transformations, we say that the actegory is strict.*

Remark 2. Although the requirement for strictness is somewhat restrictive, we will proceed under the assumption that the actegories we encounter are strict to streamline notation.

We will also be interested in actegories that interact with the monoidal structure of the underlying category.

Definition 3 (Monoidal actegory). *Let* $(\mathcal{M}, I, \otimes)$ *be a strict symmetric monoidal category and let* $(\mathcal{C}, \bullet, \eta, \mu)$ *be a strict actegory. Suppose* $\mathcal{C}$ *has a monoidal structure* $(J, \boxtimes)$*. Then we say that* $(\mathcal{C}, \bullet)$ *is monoidal if the underlying functor* $\bullet$ *is monoidal.*

We may also be interested in studying the interaction between actegorical structures and endofunctors. This interaction can happen owing to a natural transformation known as strength. We will not provide coherence diagrams in the definition below for the sake of brevity, but the interested reader can find more detail in [GLD$^+$24]. The paper also provides a definition of actegorical strong monad, which is a very similar concept.

Definition 4 (Actegorical strong functor). *Let* $(\mathcal{C}, \bullet)$ *be an* $\mathcal{M}$*-actegory. A strong actegorical endofunctor on* $(\mathcal{C}, \bullet)$ *is a pair* (F, σ) *where* $F : \mathcal{C} \to \mathcal{C}$ *is an endofunctor and* σ *is a natural transformation with components* $\sigma_{P,A} : P \bullet F(A) \to F(P \bullet A)$ *which satisfies a few coherence conditions that we do not list here.*

1.1.2 The Para construction

Suppose we have an $\mathcal{M}$-actegory $(\mathcal{C}, \bullet)$. We wish to study maps in $\mathcal{C}$ which are parametrized using objects of $\mathcal{M}$, i.e., maps in the form $P \bullet A \to B$. We are not just interested in the maps by themselves, but also in their compositional structure. Thus, we abstract away the details by defining a new category $\mathbf{Para}_\bullet(\mathcal{C})$ (first introduced in simplified form in [FST19]). Since we also want to formalize the role of reparametrization, we actually construct $\mathbf{Para}_\bullet(\mathcal{C})$ as a bicategory, so that its 0-cells A can serve as input/output spaces, its 1-cells (P, f) can serve as parametric maps, and, finally, its 2-cells r can serve as reparametrizations.

Definition 5 (**Para$_\bullet$($\mathcal{C}$)**)**.** *Let* $(\mathcal{C}, \bullet)$ *be an* $\mathcal{M}$-*actegory. Then, we define* **Para$_\bullet$($\mathcal{C}$)** *as the bicategory whose components are as follows.*

- *The* 0-*cells are the objects of* $\mathcal{C}$.
- *The* 1-*cells are pairs* $(P, f) : A \to B$, *where* $P : \mathcal{C}$ *and* $f : P \bullet A \to B$.
- *The* 2-*cells come in the form* $r : (P, f) \Rightarrow (Q, g)$, *where* $r : P \to Q$ *is a morphism in* $\mathcal{C}$. r *must also satisfy a naturality condition.*
- *The* 1-*cell composition law is*

$$(P, f) \mathbin{\mathring{,}} (Q, g) = (Q \otimes P, (Q \bullet f) \mathbin{\mathring{,}} g).$$

- *The horizontal and vertical* 2-*cell composition laws are respectively given by parallel and sequential composition in* $\mathcal{M}$.

It is quite handy to represent such cells using the string diagram notation illustrated in figure ADDIM. The **Para** *construction has a dual* **coPara** *construction whose* 1-*cells* $f : $ **coPara$_\bullet$($\mathcal{C}$)**(A, B) *take the form* (P, f), *where* $f : A \to P \bullet B$. *Cells in* **coPara$_\bullet$($\mathcal{C}$)** *can also be represented with appropriate string diagrams (see ADDIM). The reader can find a complete definition in [Gav24].*

It is shown in [Gav24] that **Para$_\bullet$($\mathcal{C}$)** is actually a 2-category if the underlying actegory is strict. Assuming this is the case (as we do in this thesis), we can use a functor $F : $ **Cat** $\to$ **Set** to quotient out the 2-categorical structure and turn **Para$_\bullet$($\mathcal{C}$)** into a 1-category $F_*($**Para$_\bullet$($\mathcal{C}$)**$)$. Here, $F_* : $ **2Cat** $\to$ **Set** is the change of enrichment basis functor induced by F. This meaningfully recovers the 1-categorical perspective of [FST19].

Both **Para$_\bullet$($\mathcal{C}$)** and **coPara$_\bullet$($\mathcal{C}$)** can be given a monoidal structure if $(\mathcal{C}, \bullet)$ is a monoidal actegory. This is extremely important because it allows us to compose (co)parametric morphisms both in sequence and in parallel. Once again, more detail can be found in [Gav24].

Remark 6. Another way to parametrize morphisms is the coKleinsli construction. As noted by [Gav24], the main difference between **coKl** and **Para** is that the parametrization offered by **coKl** is global, while the parametrization offered by **Para** is local: all morphisms in **coKl**$(X \times -)$ must take a parameter in X, while the parameter space of different morphisms of **Para**$(\mathcal{C})$ admit different parameter spaces. Nevertheless, the two constructions are related, and the former can be embedded into the latter.

If we take a parametrized category **Para$_\bullet$($\mathcal{C}$)** and we restrict our attention to morphisms parametrized with the monoidal identity I, we get back the original category $\mathcal{C}$. This is expressed by the following proposition ([Gav24]).

Proposition 7. *Let* $(\mathcal{C}, \bullet)$ *be an* $\mathcal{M}$-*actegory. Then, there exists an identity-on-objects pseudofunctor* $\gamma : \mathcal{C} \to$ **Para$_\bullet$($\mathcal{C}$)** *that maps* $f \mapsto (I, f)$. *If* $\mathcal{M}$ *is strict, this is a* 2-*functor.*

1.1.3 Optics

Modelling bidirectional flows of information is not only useful in machine learning, but also in game theory, database theory, and more. As such, categorical tools for bidirectionality have been

sought after for a long time: in particular, the greatest deal of efforts has been devoted to developing lens theory. Lenses have then been generalized into optics (see e.g. [Ril18]) to subsume other tools such as prisms and traversals into the same framework. Finally, there have also been various attempts at generalizing optics (see e.g. [CEG$^+$24] for a definition of mixed optics). We will introduce lenses and optics, and focus on the generalization of optics that appears (to us) to be the most versatile: weighted optics (first introduced in [Gav24]).

As stated in [Gav24], there is no standard definition of lens, and different authors opt for different *ad hoc* definitions that best suit their purposes. We will borrow the perspective of [CGG$^+$22] and give the following definition.

Definition 8 (Lenses). *Let $\mathcal{C}$ be a Cartesian category. Then, $\mathbf{Lens}(\mathcal{C})$ is the category defined by the following data:*

- *an object of $\mathbf{Lens}(\mathcal{C})$ is a pair $\left(\begin{smallmatrix} A \\ A' \end{smallmatrix} \right)$ of objects in $\mathcal{C}$;*
- *a $\left(\begin{smallmatrix} A \\ A' \end{smallmatrix} \right) \to \left(\begin{smallmatrix} B \\ B' \end{smallmatrix} \right)$ morphism (or lens) is a pair $\left(\begin{smallmatrix} f \\ f' \end{smallmatrix} \right)$ of morphisms of $\mathcal{C}$ such that $f : A \to B$ and $f' : A \times B' \to A'$. f is known as the forward pass of the lens $\left(\begin{smallmatrix} f \\ f' \end{smallmatrix} \right)$, whereas f' is known as the bacward pass;*
- *given $\left(\begin{smallmatrix} A \\ A' \end{smallmatrix} \right) : \mathbf{Lens}(\mathcal{C})$, the associated identity lens is $\left(\begin{smallmatrix} 1_A \\ \pi_1 \end{smallmatrix} \right)$;*
- *the composition of $\left(\begin{smallmatrix} f \\ f' \end{smallmatrix} \right) : \left(\begin{smallmatrix} A \\ A' \end{smallmatrix} \right) \to \left(\begin{smallmatrix} B \\ B' \end{smallmatrix} \right)$ and $\left(\begin{smallmatrix} g \\ g' \end{smallmatrix} \right) : \left(\begin{smallmatrix} B \\ B' \end{smallmatrix} \right) \to \left(\begin{smallmatrix} C \\ C' \end{smallmatrix} \right)$ is*

$$\left(\begin{smallmatrix} f \fatsemi g \\ \langle \pi_0, \langle \pi_0 \fatsemi f, \pi_1 \rangle \fatsemi g' \rangle \fatsemi f' \end{smallmatrix} \right) .$$

Lenses are best thought of in the helpful language of the string diagrams illustrated in ADDIM.

Lenses are a powerful tool, but they cannot be used to model all situations: for instance, lenses cannot be used if we wish to be able to choose not to interact with the environment depending on the input, or if we wish like to reuse values computed in the forward pass for further computation in the backward pass.

Optics generalize lenses by weakening the link between forward and bavkward passes, and by replacing the Cartesian structure of the underlying category with a simpler symmetric monoidal structure. In an optic over $\mathcal{C}$, an object $M : \mathcal{C}$ acts as an inaccessible residual space transferring information between the upper components and the lower component. We will provide the definition given by [Ril18][1], but we will use the string diagram notation presented by [Gav24] for the sake of consistency.

Definition 9 (Optics). *Let $(\mathcal{C}, I, \otimes, \lambda, \rho, \alpha)$ be a symmetric monoidal category (we make the unitors and associators explicit for later use). Then, $\mathbf{Optic}(\mathcal{C})$ is the category defined by the following data:*

- *an object of $\mathbf{Optic}(\mathcal{C})$ is a pair $\left(\begin{smallmatrix} A \\ A' \end{smallmatrix} \right)$ of objects in $\mathcal{C}$;*

---

[1][Ril18]also provides a more versatile (but more sophisticated) definition of optics that relies on coends. Under the coend formalism,

$$\mathbf{Optic}(\mathcal{C}) \left(\left(\begin{smallmatrix} A \\ A' \end{smallmatrix} \right), \left(\begin{smallmatrix} B \\ B' \end{smallmatrix} \right) \right) = \int^{M : \mathcal{C}} \mathcal{C}(A, M \times B) \times \mathcal{C}(M \otimes B', A').$$

- a $\left(\begin{smallmatrix} A \\ A' \end{smallmatrix} \right) \to \left(\begin{smallmatrix} B \\ B' \end{smallmatrix} \right)$ morphism (or optic) is a pair $\left(\begin{smallmatrix} f \\ f' \end{smallmatrix} \right)$ of morphisms of $\mathcal{C}$ such that $f : A \to M \otimes B$ and $f' : M \otimes B' \to A'$, where $M : \mathcal{C}$ is known as residual space; such pairs $\left(\begin{smallmatrix} f \\ f' \end{smallmatrix} \right)$ are also quotiented by an equivalence relation that allows for reparametrization of the residual space and effectively makes it inaccessible;
- the identity on $\left(\begin{smallmatrix} A \\ A' \end{smallmatrix} \right)$ is the optic represented by $(\lambda_A^{-1}, \lambda_A)$.

We will only show how optics compose using the string diagrams in ADDIM (see [Ril18] for the composition formula).

Lenses come up as a special case of optics ([Ril18]), and optics do solve some of the issues we have with lenses. However, optics are not perfect either: for instance, [Gav24] points out that optics cannot be used in cases where we ask that the forward pass and backward pass are different kind of maps, as they are both forced to live in the same category. Thus, we need a further layer of generalization: namely, weighted optics.

1.1.4   Weighted optics

Before we define weighted optics, we need to introduce a new tool to our toolbox: the category of elements of a functor.

Definition 10 (Elements of a functor)**.** *Let $F : \mathcal{C} \to \mathbf{Set}$ be a functor. We define $\mathbf{El}(F)$ as the category with the following data: (i) the objects of $\mathbf{El}(F)$ are pairs (C, x) where $C : \mathcal{C}$ and $x : F(C)$; (ii) the $(C, x) \to (D, y)$ morphisms in $\mathbf{El}(F)$ are the morphisms $f : C \to D$ in $\mathcal{C}$ such that $F(f)(x) = y$.*

[Gav24] studies $\mathcal{B}$-actegories $(\mathcal{C}, \bullet)$, which are then reparametrized so that the acting category becomes $\mathcal{E} = \mathbf{El}(W)$ for some weight functor $W : \mathcal{B}^{\mathrm{op}} \to \mathbf{Set}$ (which is to be specified). The reparametrization takes place thanks to the opposite of the forgetful functor $\pi_W : \mathcal{E} \to \mathcal{B}^{\mathrm{op}}$, which maps $(B, x) \mapsto B$. Hence, we consider the action

$$\bullet^{\pi_W^{\mathrm{op}}} = \mathbf{El}(W)^{\mathrm{op}} \times \mathcal{C} \xrightarrow{\pi_W^{\mathrm{op}} \times \mathcal{C}} \mathcal{B} \times \mathcal{C} \xrightarrow{\bullet} \mathcal{C}.$$

We are finally ready to define weighted optics[2].

Definition 11 (Weighted **coPara**)**.** *If W is a weight functor as above and $(\mathcal{C}, \bullet)$ is a $\mathcal{B}$-actegory, we define*

$$\mathbf{coPara}_\bullet^W(\mathcal{C}) = \pi_{0*}(\mathbf{coPara}_{\bullet^{\pi_W^{\mathrm{op}}}}(\mathcal{C})),$$

where π_{0*} is the enrichment base change functor generated by the connected component functor $\pi_0 : \mathbf{Cat} \to \mathbf{Set}$. More explicitly, π_{0*} quotients the connections provided by reparametrizations.

Definition 12 (Weighted optics)**.** *Suppose $(\mathcal{C}, \bullet)$ is an $\mathcal{M}$-actegory, suppose $(\mathcal{D}, \blacksquare)$ is an $\mathcal{M}'$-actegory, and suppose $W : \mathcal{M}^{\mathrm{op}} \times \mathcal{M}' \to \mathbf{Set}$ is a lax monoidal functor. We define the category of W-weighted optics over the product actegory $(\mathcal{C} \times \mathcal{D}^{\mathrm{op}}, (\, {}_{\blacksquare^{\mathrm{op}}}^\bullet \,))$ as*

$$\mathbf{Optic}_{\left(\begin{smallmatrix} \bullet \\ \blacksquare \end{smallmatrix} \right)}^W = \mathbf{coPara}_{\left(\begin{smallmatrix} \bullet \\ \blacksquare^{\mathrm{op}} \end{smallmatrix} \right)}^W(\mathcal{C} \times \mathcal{D}^{\mathrm{op}}).$$

---

[2]Weighted optics also admit a coend definition. Refer to [Gav24] for more information.

The definition is very dense and deserves some explanation. Fist of all, we assume that W maps (M, M') to a set of maps $s : M \to M'$. If that's the case, a $\begin{pmatrix} X \\ X' \end{pmatrix} \to \begin{pmatrix} Y \\ Y' \end{pmatrix}$ map is a triplet $\left((M'), s, \begin{pmatrix} f \\ f' \end{pmatrix} \right)$, where M is the forward residual, M' is the backward residual, $s : M \to M'$ links the two residuals, $f : X \to M \bullet Y$ is the forward pass, and $f' : M' \bullet Y' \to X'$ is the backward pass. The triplets are also quotiented with respect to reparametrization, which makes the residual spaces effectively inaccessible (as it happens in the case of ordinary optics). We can get a clear "operational" understanding of how a weighted optic works looking at an associated string diagram: data from X flows through the forward map, which computes an output in Y and a forward residual in M. Such forward residual is then converted into a backward residual in M' by the map s, which is provided by the weight functor. Finally, the backward residual is used to compute, together with input from Y' a value in X'. The last step happens thanks to the backward map f'. A full account of the composition law for weighted optics can be found on [Gav24]. As stated in [Gav24], since **coPara** can be given a monoidal structure, we can also give $\textbf{Optic}^{W}_{\begin{pmatrix} \bullet \\ \bullet \end{pmatrix}}$ one such structure as long as the underlying actegories are monoidal and the weight functor W is braided monoidal.

The advantages of weighted optics over ordinary optics are clear: when dealing with weighted optics, we are no longer forced to take reverse maps from the same category as the forward maps. The action on the category of forward spaces is now separated from the action on the category of backward spaces, and the link between the two actions is provided by an external functor. Such modular approach provides a great deal of conceptual clarity and flexibility, more than regular optics or lenses can provide on their own. It is also shown in [Gav24] that weighted optics are indeed a generalization of optics. In particular, it is shown that the lenses in *Def. 8* are the specialized weighted optics obtained when $\mathcal{C} = \mathcal{D}$ is Cartesian and the actegories are given by the Cartesian product. More generally, [Gav24] claims that - to the best of the author's knowledge - all definitions of lenses currently used in the literature are subsumed by the definition of weighted optics.

[Gav24] goes on to apply the **Para** construction onto weighted optics, obtaining parametric weighted optics, which are proposed as a full-featured model for deep learning. The author conjectures that "weighted optics provide a good denotational and operational semantics for differentiation". In its full, generality, this is still an unproven conjecture. However, restricting our attention to a special class $\textbf{Lens}_A$ of lenses with an additive backward passes yields a fully formal theory of structural back-propagation, which will be illustrated in the rest of the capter, after a short digression on differential categories.

1.1.5   Differential categories

Modelling gradient-based learning obviously requires a setting where differentiation can take place. Although it is tempting to directly employ smooth functions over Euclidean spaces, recent research has shown that there are tangible advantages in working with generalized differential combinators that extend the notion of derivative to polynomial circuits ([WZ22], [WZ21]), manifolds (ADDREF), complex spaces (ADDREF), and so on. Thus, it makes sense to work with an abstract notion of derivative which can then be appropriately implemented depending on the requirements at hand.

The standard approach to this problem (developed in the last twenty years) involves the explicit definition of two kidns of differential categories: Cartesian differential categories (first introduced in [BCS06]) and Cartesian reverse differential categories (first introduced by [CCG$^+$19]). The former

allow for forward differentiation, while the latter allow for reverse differentiation. We will omit the defining axioms for the sake of brevity, but the reader can find complete definitions in [CCG$^+$19].

Definition 13 (Cartesian differential category)**.** *A Cartesian differential category (CDR) $\mathcal{C}$ is a Cartesian left-additive category where a differential combinator D is defined. Such differential combinator must take a morphism $f : A \to B$ and return a morphism $\mathrm{D}[f] : A \times A \to B$, which is known as the derivative of f. The combinator D must satisfy a number of axioms.*

Definition 14 (Cartesian reverse differential category)**.** *A Cartesian reverse differential category (CRDC) $\mathcal{C}$ is a Cartesian left-additive category where a reverse differential combinator R is defined. Such reverse differential combinator must take a morphism $f : A \to B$ and return a morphism $\mathrm{R}[f] : A \times B \to A$, which is known as the reverse derivative of f. The combinator R must satisfy a number of axioms.*

Example 15. Smooth is a both a CDC and a CRDC. In fact, if $\mathcal{J}_f$ is the Jacobian matrix of a smooth morphism f,

$$\mathrm{D}[f] : (x, v) \mapsto \mathcal{J}_f(x)v$$

and

$$\mathrm{R}[f] : (x, y) \mapsto \mathcal{J}_f(x)^T y$$

induce well-defined combinators D and R. This is only a partial coincidence, as shown in [CCG$^+$19] that CRDCs are always CDCs under a canonical choice of differential combinator. The converse, however, is generally false.

As it turns out, forward differentiation tends to be less efficient when dealing with neural networks that come up in practice (ADDREF), so CDCs are not extremely useful when studying deep learning. CRDCs, on the other hand, have been applied to great success (see e.g. [CGG$^+$22]). As shown in [WZ22], a large supply of CRDCs can be obtained by providing the generators of a finitely presented Cartesian left-additive category with associated reverse derivatives (as long as the choices of reverse derivative are consistent). Moreover, CRDCs have been recently generalized by [Gav24] to coalgebras associated with copointed endofunctors, which could also increase the number of known CRDCs in the future. The rest if this section is devoted to this generalization.

It is shown in [Gav24] that there is a particular class of weighted optics which is useful for reverse differentiation, being able to represent both maps (through forward passes) and the associated reverse derivatives (through backward passes). Moreover, such weighted optics can be represented as lenses in the sense of *Def 8*, which means that their inner workings can be pictured in a simple, intuitive way.

Definition 16 (Additively closed Cartesian left-additive category)**.** *A Cartesian left-additive category $\mathcal{C}$ is an additively closed Cartesian left-additive category (ACCLAC) if and only if the following are true:*

- *the subcategory $\mathbf{CMon}(\mathcal{C})$ of additive maps has a closed monoidal structure $(I, \otimes)$;*
- *the embedding $\iota : \mathbf{CMon}(\mathcal{C}) \to \mathcal{C}$ is a lax monoidal funtor with respect to the aforementioned structure of $\mathbf{CMon}(\mathcal{C})$ and the Cartesian structure of $\mathcal{C}$.*

Then, we can define the category of lenses with backward passes additive in the second component.

Definition 17. *Let $\mathcal{C}$ be an ACCLAC with Cartesian structure is $(1, \times)$ and whose subcategory* $\mathbf{CMon}(\mathcal{C})$ *has monoidal structure $(I, \otimes)$. Then, we define*

$$\mathbf{Lens}_A(\mathcal{C}) = \mathbf{Optic}_{\binom{\times}{\otimes}}^{\mathcal{C}(-, \iota(-))}.$$

As argued in [Gav24], the symbol $\mathbf{Lens}_A$ is justified because one such optic of type $\binom{X}{X'} \to \binom{Y}{Y'}$ can be concretely represented as a lens with forward pass $f : \mathcal{C}(X, Y)$ and backward pass $f' : \mathcal{C}(X \times Y', X')$, which is the approach we illustrate in this thesis. Nevertheless, some potential expressiveness is lost when passing from weighted optic composition to concrete lens composition. In particular, if we operated with optics, we would be able to implement backpropagation without resorting to gradient checkpointing, which is not possible if we use lenses ([Gav24]).

The generalization mentioned above is possible because $\mathbf{Lens}_A$ is an endofunctor.

Definition 18. *We defined* $\mathbf{CLACat}$ *as the category whose objects are Cartesian left-additive categories and whose morphisms are Cartesian left-additive functors (see e.g.[BCS06]).*

Proposition 19. *If $\mathcal{C} : \mathbf{CLACat}$, then $\mathbf{Lens}_A(\mathcal{C}) : \mathbf{CLACat}$.*

Proof. The Cartesian structure on $\mathbf{Lens}_A(\mathcal{C})$ is given by $\binom{X}{X'} \times \binom{Y}{Y'} = \binom{X \times Y}{X' \times Y'}$ and by the initial object $\binom{1}{1}$. The monoidal structure on each $\binom{X}{X'}$ is given by the unit $0_{\binom{X}{X'}} = \binom{0_A}{!_{1 \times A'}} : \binom{1}{1} \to \binom{X}{X'}$ and by the multiplication $+_{\binom{X}{X'}} = \binom{+_A}{\pi_2 \, \mathring{,} \, \Delta_{A'}} : \binom{X \times X}{X' \times X'} \to \binom{X}{X'}$. $\square$

Proposition 20. $\mathbf{Lens}_A : \mathbf{CLACat} \to \mathbf{CLACat}$ *is a functor.*

Proof. Given a Cartesian left-additive functor $F : \mathcal{C} \to \mathcal{D}$, we can define $\mathbf{Lens}_A(F)$ as the functor that maps $\binom{X}{X'} \mapsto \binom{F(X)}{F(X')}$ and maps $\binom{f}{f'} \mapsto \binom{F(f)}{\underline{f}}$, where $\underline{f} = F(X) \times F(Y')) \xrightarrow{\cong} F(X \times Y') \xrightarrow{F(f')} F(X')$. It can be shown that $\mathbf{Lens}_A(F)$ is also Cartesian left-additive. $\square$

Proposition 21. $\mathbf{Lens}_A$ *has a copointed structure*[3].

Proof. It suffices to endow $\mathbf{Lens}_A$ with the natural transformation ϵ whose components are the forgetful functors $\epsilon_{\mathcal{C}} : \mathbf{Lens}_A(\mathcal{C}) \to \mathcal{C}$ which strip away the backward passes. $\square$

Hence, [Gav24] defines generalized CRDCs as follows.

Definition 22 (Generalized Cartesian reverse differential category)**.** *A generalized Cartesian reverse differential category is a coalgebra for the pointed endofunctor* $\mathbf{Lens}_A$.

Explicitly, a colagebra for $\mathbf{Lens}_A$ is a pair $(\mathcal{C}, \mathbf{R}_{\mathcal{C}})$ such that $\mathcal{C} : \mathbf{CLACat}$ and $\mathbf{R}_{\mathcal{C}} : \mathcal{C} \to \mathbf{Lens}_A(\mathcal{C})$ satisfies $\mathbf{R}_{\mathcal{C}} \, \mathring{,} \, \epsilon_{\mathcal{C}} = \mathrm{id}_{\mathcal{C}}$. The intuition behind such definition is that $\mathbf{R}_{\mathcal{C}}$ should map $f \mapsto \binom{f}{R[f]}$, where $R[f]$ is a generalized reverse derivative combinator. [Gav24] shows that such a definition of $\mathbf{R}_{\mathcal{C}}$ does indeed prove that ordinary CRDC fall into the definition of generalized CRDC.

---

[3]An endofunctor $F : \mathcal{C} \to \mathcal{C}$ is copointed if it is endowed with a natural transformation $\epsilon : F(\mathcal{C}) \Rightarrow \mathrm{id}_{\mathcal{C}}$.

1.1.6   Parametric lenses

We conclude this section discussing the relation between the **Para** construction and the $\mathbf{Lens}_A$ endofunctions. [Gav24] shows that, under an appropriate definition, 'morphisms of actegories induce morphisms of parametric bicategories'. As a consequence, it can be shown that, if $(\mathcal{C}, \mathbf{R}_\mathcal{C})$ is a generalized CRDC, $\mathbf{R}_\mathcal{C}$ induces a functor $\mathbf{Para}(\mathbf{R}_\mathcal{C}) : \mathbf{Para}_\times(\mathcal{C}) \to \mathbf{Para}_\bullet(\mathbf{Lens}_A(\mathcal{C}))$, which takes a parametric map $f : P \times A \to B$ and augments is with its reverse derivative $\mathrm{R}[f]$, forming a parametric lens. Parametric lenses behave very similarly to lenses, but we provide a separate stand-alone definition (which we take from [CGG$^+$22]) for the reader's convenience.

Definition 23 (Parametric lenses)**.** *The category of parametric lenses over a Cartesian category* $(\mathcal{C}, 1, \times)$ *is* $\mathbf{Para}_\bullet(\mathbf{Lens}(\mathcal{C}))$, *where $\bullet$ is the action on the lenses generated by the Cartesian structure of $\mathcal{C}$:*

$$\begin{pmatrix} P \\ P' \end{pmatrix} \bullet \begin{pmatrix} A \\ A' \end{pmatrix} = \begin{pmatrix} P \times A \\ P' \times A' \end{pmatrix}.$$

Refer to Fig. ADDIM to see a string diagram that shows the inner workings of a parametric lens.

1.2   Supervised learning with parametric lenses

In this section, we show how parametric lenses can be used to model supervised gradient-based learning ([CGG$^+$22], [Gav24], [SGW21]). While lenses are not as general as weighted optics, it is shown in [CGG$^+$22] that they are powerful enough for most purposes and that there is some empirical evidence of their performance and applicability. The paper also discusses the use of parametric lenses in modeling unsupervised deep learning and deep dreaming, but we do not have the space to discuss them here.

1.2.1   Model, loss, optimizer, learning rate

Supervised gradient-based learning can be modeled using parametric lenses as follows:

1. we can design an architecture (P, Model) as a parametric morphism in $\mathbf{Para}_\bullet(\mathcal{C})$ for some generalized CRDC $(\mathcal{C}, \mathbf{R}_\mathcal{C})$;
2. we can use the functor $\mathbf{R}_\mathcal{C}$ to endow (P, Model) with its reverse derivative $\mathrm{R}[(P, \mathsf{Model})]$, yielding a lens in $\mathbf{Para}_\bullet(\mathbf{Lens}_A(\mathcal{C}))$;
3. we can use 2-categorical machinery of $\mathbf{Para}_\bullet(\mathbf{Lens}_A(\mathcal{C}))$ to provide a loss function, a learning rate, and optimizer, which can be assembled onto $\mathbf{R}_\mathcal{C}(P, \mathsf{Model})$ to yield a parameter update;
4. we can use copy maps from the Cartesian structure of $\mathcal{C}$ to create a learning iteration.

The theory of parametric optics and differential categories does not offer explicit insight with respect to architecture design, so we will assume a good architecture as already been designed (at least for now). Given an architecture (P, Model), it can be embedded into $\mathbf{Para}_\bullet(\mathbf{Lens}_A(\mathcal{C}))$ as a lens $\begin{pmatrix} P \times A \\ P \times A \end{pmatrix} \to \begin{pmatrix} B \\ B \end{pmatrix}$ by breaking it up into its basic components (such as linear layers, convolutional layers, etc.), augmenting such components with their reverse derivatives, and the composing the resulting

lenses. The backward pass of the composition is the reverse derivative of its forward pass because $\mathbf{R}_{\mathcal{C}}$ is a functor[4]. Many examples can be found in [CGG+22].

Updating the parameters based on data requires a loss function, an optimizer and a learning rate. Loss functions can be implemented as parametric lenses which take in predictions as input and labels as parameters. The output they produce can be considered the actual loss that needs to be differentiatied. Given a model parametric lens $\left(\begin{smallmatrix} \mathsf{Model} \\ \mathsf{R[Model]} \end{smallmatrix}\right) : \left(\begin{smallmatrix} P \times A \\ P \times A \end{smallmatrix}\right) \to \left(\begin{smallmatrix} B \\ B \end{smallmatrix}\right)$ and a loss parametric lens $\left(\begin{smallmatrix} \mathsf{Loss} \\ \mathsf{R[Loss]} \end{smallmatrix}\right) : \left(\begin{smallmatrix} B \times B \\ B \times B \end{smallmatrix}\right) \to \left(\begin{smallmatrix} L \\ L \end{smallmatrix}\right)$, the composition $m \, \mathring{,} \, l$ takes in features as input and takes model parameters and labels as parameters. Then, $\left(\begin{smallmatrix} \mathsf{Model} \\ \mathsf{R[Model]} \end{smallmatrix}\right) \, \mathring{,} \, \left(\begin{smallmatrix} \mathsf{Loss} \\ \mathsf{R[Loss]} \end{smallmatrix}\right)$ computes the loss associated with the model predictions. See ADDIM for the associated string diagram.

It can be helpful to think about dangling wires in the diagrams as open slots where other components can be plugged. For instance, the diagram of ADDIM has dangling wires labeled with L on its right. We can use a learning rate lens α to link these wires and allow forward-propagating information to "change direction" and go backwards. α must have domain equal to $\left(\begin{smallmatrix} L \\ L \end{smallmatrix}\right)$ and codomain equal to $\left(\begin{smallmatrix} 1 \\ 1 \end{smallmatrix}\right)$, where 1 is the terminal object of $\mathcal{C}$. For instance, if $\mathcal{C} = \mathbf{Smooth}$, α might just multiply the loss by some ϵ, which is what machine learning practitioners would ordinarily call learning rate. Fig. ADDIM shows how a learning rate can be linked to the loss function and the model using post-composition.

The final element needed for the model Model in Fig. ADDIM to learn is an optimizer. It is shown in [CGG+22] that optimizers can be represented as reparametrisations in $\mathbf{Para}(\mathbf{Lens}(\mathcal{C}))$. More specifically, we might see an optimizer as a lens $\left(\begin{smallmatrix} P \\ P \end{smallmatrix}\right) \to \left(\begin{smallmatrix} Q \\ Q \end{smallmatrix}\right)$. In gradient descent, for example, $P = Q$ and the aforementioned lens is $\left(\begin{smallmatrix} 1_P \\ +_P \end{smallmatrix}\right)$. We can plug such reparametrisation on top of the model to obtain the string diagram in Fig. ADDIM. The diagram shows how the machinery hidden by the $\mathbf{Para}(\mathbf{Lens}(\mathcal{C}))$ can take care of forward propagation, loss computation, backpropagation and parameter updating in a seamless fashion.

The supervised learning lens in ADDIM actually has two more dangling wires we don't have any use for: namely, the second input A of the model and the second parameter B of the loss functions. These are not a problem as we can just plug them with delete maps taken from the Cartesian structure of $\mathcal{C}$. We can also bend the input wires with an appropriate lens to turn them into parameters. We are then left with a $\left(\begin{smallmatrix} 1 \\ 1 \end{smallmatrix}\right) \to \left(\begin{smallmatrix} 1 \\ 1 \end{smallmatrix}\right)$ parametric lens with parameter space $\left(\begin{smallmatrix} A \\ 1 \end{smallmatrix}\right) \times \left(\begin{smallmatrix} P \\ P \end{smallmatrix}\right) \times \left(\begin{smallmatrix} B \\ 1 \end{smallmatrix}\right)$. This lens is pictured in ADDIM.

1.2.2 Weight tying, batching, and the learning iteration

Both [CGG+22] and [Gav24] emphasize the essential role played by weight tying in deep learning. Weight tying can be implemented within the parametric lens framework as a reparametrization that copies a single parameter to many parameter slots (see ADDIM): given $(P \times P, f) : \mathbf{Para}(\mathcal{C})(X, Y)$, we can define $(P, f^{\Delta_P}) : \mathbf{Para}(\mathcal{C})(X, Y)$ so that

$$f^{\Delta_P} : P \times X \xrightarrow{\Delta_P \times X} P \times P \times X \xrightarrow{f} Y.$$

Weight tying can also be used for batching: batching is implemented by instantiating n different copies of our supervised learning lens (comprised of model, loss function, and learning rate) and tying the

---

[4]As highlighted by [SGW21], the diagram for the backward pass of the composition of two lenses looks exactly like the diagram describing the chain rule for reverse derivatives, which is what makes $\mathbf{R}_{\mathcal{C}}$ a well-defined functor.

parameters to a unique value. Then, it suffices to feed the n data points to the n lenses, and we can optimize across a single parameter (see ADDIM).

[CGG$^+$22] introduces a possible representation for the whole learning iteration of a supervised learning model within the parametric lens framework. The paper suggests extracting the bacward pass of the lens in ADDIM and reframing it as a $P \to P$ parametric map with parameters $A \times B$. Since this is an endomap, it can be composed n times with itself to obtain a $P \to P$ map, which is proposed as a model of the learning iteration. While this approach requires breaking lenses apart, it is markedly simple.

1.2.3   Empirical evidence

Empirical evidence for the effectiveness of the parametric lens framwork discussed in this section can be found in [CGG$^+$22], where the authors go on to implement a Python libraries for gradient-based learning based on these ideas. They use the library to develop a MNIST classifier, obtaining comparable accuracy to models developed using traditional tools.

Since parametric lenses can easily be implemented functionally and without side effects, success stories as the one mentioned above foreshadow a future where popular machine learning libraries also follow elegant functional paradigms informed by category theory. Quoting [CGG$^+$22] directly, '[the] proposed algebraic structures naturally guide programming practice'.

1.3   Future directions and related work

The parametric optic framework discussed in this chapter is very promising, but there is still a lot of work that needs to be done so that it can reach its full potential. For instance, [Gav24] conjectures that weighted optics can be used in its full generality to model differentiation in cases which are not covered by lenses. For instance, lenses cannot model AD algorithms that do not use gradient checkpointing, while weighted optics are conjectured to be able to do so. [Gav24] also suggests investigating locally graded categories as potential replacements for actegories, and also investigating the applications of parametric optics to meta-learning, i.e. deep learning where the optimizers themselves are learned. Moreover, [CGG$^+$22] conjectures that some of the axioms of CRDC may be used to model higher order optimization algorithms. Finally, as suggested by [CGG$^+$22], hopefully future work will allow the parametric optic framework to encompass non-gradient based optimizers such as the ones used in probabilistic learning. See [SGW21] for more on this topic.

We conclude this chapter by discussing three other directions of machine learning reasearch that are closely related to the framework of parametric optics.

1.3.1   Learners

One of the first compositional approaches to training neural networks in the literature can be found in the seminal paper [FST19], which spurred much research in the field, including what is presented in [Gav24] and [CGG$^+$22]. The authors introduce a category of learners, objects which are meant to represent components of a neural network and behave similarly to parametric lenses.

Definition 24 (Category of learners)**.** *Let A and B be sets. A learner $A \to B$ is a tuple (P, I, U, r) where P is a set, and $I : P \times A \to B$, $U : P \times A \times B \to P$, and $r : P \times A \times B \to A$ are functions. P is known as parameter space, I as implement functions, U as update function, and r as request function. Two learners $(P, I, U, r) : A \to B$ and $(Q, J, V, s) : B \to C$ compose forming $(P \times Q, I * J, U * V, r * s) : A \to C$, where*

$$(I * J)(p, q, a) = J(q, I(p, a)),$$

$$(U * V)(p, q, a, c) = (U(p, a, s(q, I(p, a), c)), V(q, I(p, a), c)),$$

$$(r * s)(p, q, a, c) = r(p, a, s(q, I(p, a), c)).$$

Learners quotiented by an appropriate reparametrization relationship[5] form a category **Learn***.*

A learner represents an instance of supervised learning: the implement function takes a parameter and implements a function and the update function updates the parameters using a data from a dataset. The request function is necessary to implement backpropagation when optimizing a composition of learners. Suppose we select a learning rate ϵ and an error function $e : \mathbb{R}^2 \to \mathbb{R}$ such that $y \mapsto \frac{\partial e}{\partial x}(x_0, y)$ is invertible for all y. It is argued in [FST19] that we can define a functor $L_{\epsilon,e} : \mathbf{Para}_\times(\mathbf{Smooth}) \to \mathbf{Learn}$ which takes a parametric map and yields an associated learner that implements gradient descent.

We do not have the space to talk about learners at length, but we wish to draw a short comparison between parametric weighted optics (and, in particular, parametric lenses) and the approach of [FST19], given the relevant position held by the paper in the machine learning literature. The similarities between learner-based learning and lens-based learning are evident: every learner (P, I, U, r) looks like a parametric lens, where I passes information forward, r passes information backwards and P is the parameter space. Moreover, the role of $L_{\epsilon,e}$ is very similar to the role played by $\mathbf{Para}(\mathbf{R}_\mathcal{C})$ in optic-based learning. Such similarities were even discussed in the original paper [FST19] and have been researched at length: it has been proved in [FJ19] that learners can be functorially and faithfully embedded in a special category of symmetric lenses (as opposed to the lenses of *Def. 8*, which are asymmetric).

Despite the similarities, there is one fundamental difference between the lens-based approach and the learner-based approach: each learner carries its own optimizer, whereas optimization of lenses is usually carried out separately. Moreover, if we compare parametric weighted optics with learners, the latter clearly win in versatility, generality, and (at least from our point of view) conceptual clarity. It is argued in [SGW21] and [CGG$^+$22] that the parametric lens framework largely subsumes the learner approach. More information regarding the comparison can also be found in [Gav24].

1.3.2 Exotic differential categories

We have presented the parametric weighted optic approach of [Gav24] and [CGG$^+$22] within the context of neural networks for the sake of simplicity, but the framework has been developed with generality in mind and applies to a much wider range of situations. For instance, we can easily replace

---

[5]As argued in [FST19], learners could be studied from a bicategorical point of view, where reparametrizations would just be 2-cells. We could then use a connected component projection to compress **Learn** into a 1-category **Learn**, as it is done for **coPara** when defining weighted optics.

Smooth with any other CRDC $\mathcal{C}$, yielding a full-feature compositional framework for gradient-based learning over $\mathcal{C}$.

Switching to a different CRDC is useful because different differential categories can lead to very different learning outcomes, both in terms of accuracy of the model and in terms of training computational costs (ADDREF). For instance, is argued in [WZ22] that polynomial circuits can be used to define and train intrinsecally discrete machine learning models. Even 'radical 'environments such as Boolean circuits - where scalars reside in $\mathbb{Z}_2$ - seem to be conductive to machine learning under the right choice of architecture and optimizer ([WZ21]). Using such exotic differential categories may be of great advantage because they might be able to better reflect the intrinsic computational limits of computer arithmetics, leading to more efficient learning ([WZ22]).

1.3.3 Functional reverse-mode automatic differentiation

Finally, we wish to highlight the similarities between the formal theory of differential categories illustrated here and the work in [Ell18]. The paper describes the Haskell implementation of a purely functional automatic differentiation library, which is able to handle both forward mode and backward mode AD without resorting to the mutable computational graphs used by most current day libraries.

Among the main insights of [Ell18], it is stated that derivatives should not be treated as simple vectors, but as linear maps, or multilinear maps in the case of uncurried higher-order derivatives. Moreover, the author shows that differentiation can be made compositional by working on pairs (f, Df), which behaved very similarly to lenses. As noted by [SGW21], however, [CGG$^+$22] and other lens-theoretical perspectives do not subsume the work in [Ell18] because of the latter's programming focus. See [SGW21] for more information regarding this comparison.

Chapter 2

From Classical Computer Science to Neural Networks

Classical computer science focuses on discovering algorithms, i.e. ordered sequences of steps which operate in precisely set, idealized conditions and have strong guarantees of correctness due to their exact mathematical formulations. Neural networks, on the other hand, are able to work in messy, real-world conditions, but offer very so few guarantees of correctness that their performance is often described as *unreasonably* good. Moreover, whereas algorithms generalize very well (most software engineers will only need a few dozen algorithms in their entire career), neural networks are often completely helpless when pitted against out of distribution inputs. Hence, algorithms and neural networks can be seen as complementary opposites ([VB21], [VBB+22]).

Recent attempts group under then label of *neural algorithmic reasoning* (see [VB21] for a very short introduction to the subject) have tried to get the best of both worlds by training neural networks to apply algorithms (see e.g [IKP+22]). The CLRS benchmark (introduced by [VBB+22]) uses graphs to represent the computations associated with a few classical algorithms from the famous CLRS introductory textbook so that graph neural networks (GNNs) can be trained to learn these algorithms. The benchmark has spurred a large amount of research in this direction, with very promising results.

More generally, linking machine learning to classical informatics may hold the key for interesting advances. For example, recovering neural networks as parametric versions of known algorithms may help imbue the ever chaotic field of machine learning with new-found conceptual clarity. Moreover, such a link could even help develop new neural network architectures by taking inspiration from well-researched classical notions. In this chapter we will illustrate two lines of inquiry which use category theory to build such a bridge: *categorical deep learning* and an interesting categorical approach to GNNs. Before proceeding with such topics, we will go on a short tangent categorical tangent regarding (co)algebras and the integral transform.

2.1 Categorical toolkit

2.1.1 (Co)algebras

Algebras and coalgebras are a categorical formalization of the principles of induction and coinduction. Induction and coinduction are fundamental to computer science because the allow us to give precise definitions for a very large number of data structures, and because they allow us to formalize the algorithms that operate on these structure within the schemas of recursion and corecursion. We will touch on (co)algebras very briefly and we encourage the interested reader to refer to [JR97] and [Wis08] for further detail.

Definition 25 ((Co)algebra over an endofunctor). *Let $F : \mathcal{C} \to \mathcal{C}$ be an endofunctor. An algebra over F is a pair (A, a) where $A : \mathcal{C}$ and $a : \mathcal{C}(F(A), A)$. A coalgebra is a pair (A, a) where $A : \mathcal{C}$ and $a : \mathcal{C}(A, F(A))$. In both cases A is known as carrier set and a as structure map.*

(Co)algebras can also be defined on monads: the only difference (co)algebras over an endofunctor and (co)algebras over a monad is that the latter also need to be compatbile with the monad structure, i.e. satisfy commutative diagrams that represent coherence conditions (see [GLD+24]).

(Co)algebras over the same functor can be given a categorical structure by using the following notion of homomorphism.

Definition 26 (Homomorphisms of (co)algebras over an endofunctor). *Let (A, a) and (B, b) be algebras over the same endofunctor $F : \mathcal{C} \to \mathcal{C}$. An algebra homomorphism $(A, a) \to (B, b)$ is a map $f : \mathcal{C}(A, B)$ such that the diagram in ADDIM is commutative.*

Now suppose (A, a) and (B, b) are coalgebras. A homomorphism between them is a map $f : \mathcal{C}(A, B)$ such that the diagram in ADDIM is commutative.

The main intuition behind algebras and coalgebras is the following: the underlying functor defines a signature for the (co)algebraic structure; the structure of an algebra is a constructor that takes data from $F(A)$ and uses it to build data from A, whereas the structure of a coalgebra observes data from A and produces an observation in the form of data from $F(A)$; (co)algebra homomorphism are arrows that preserve the underlying structure. Consider the following clarifying examples from [GLD+24][1].

Example 27 (Lists). Let A be a set. Consider the polynomial endofunctor $1 + A \times X : \mathbf{Set} \to \mathbf{Set}$. If $\mathsf{List}(A)$ is the set of A-labeled lists, $(\mathsf{List}(A), [\mathsf{Nil}, \mathsf{Cons}])$ is an algebra over $1 + A \times X$. Here, $\mathsf{Nil} : 1 \to List(A)$ is the map which takes the unique object of 1 and returns the empty list, while $\mathsf{Cons} : A \times \mathsf{List}(A) \to \mathsf{List}(A)$ is the map which takes an element $a \in A$ and a list l of elements of A and returns the concatenated list $l \cup \{a\}$. This described lists in $\mathsf{List}(A)$ inductively as object formed by concatenating elements of A to other such lists. The base case is the empty list.

Example 28 (Mealy machines). Now consider two sets I and O of possible inputs and outputs, respectively. Consider the endofunctor $(O \times X)^I : \mathbf{Set} \to \mathbf{Set}$. Define $\mathsf{Mealy}_{I,O}$ as the set of Mealy machines with inputs and outputs in I and O, respectively. Now we can consider the coalgebra

---

[1]In the examples below we use polynomial and exponential expressions to define endofunctors over **Set**. In this context, X is the argument of the functor, $\times$ is the Cartesian product, $+$ is the disjoint union, $\langle f, g \rangle$ is the pairing induced by $\times$, $[f, g]$ is the pairing induced by $+$, and B^A is the set of functions $A \to B$.

($\mathsf{Mealy}_{I,O}$, Next), where Next is the map that takes a Mealy machine $m \in \mathsf{Mealy}_{I,O}$ and yields a function which in turn, given $i \in I$, returns the output of m at i and a new machine. This is a coinductive description of Mealy machines.

Remark 29. Notice how the description we have given of Mealy machines does not mention the internal states of these objects. This is a recurring aspect of coinductive descriptions: as argued in [JR97], coinduction is best interpreted as a process where an observer track the behavior of an object from the outside, with no access to its internal state. This is very useful in machine learning because the internal state of a learning model is often unknown or uninterpretable.

The link between (co)algebras and (co)induction does not stop at the definition level. The example below shows that an algebra homomorphism can model a recursive fold procedure. A similar corecursive unfold procedure can be defined by using a coalgebra homomorphism (see [GLD$^+$24] for further detail).

Example 30 (List folds)**.** Consider the algebra $(\mathsf{List}(A), [\mathsf{Nil}, \mathsf{Cons}])$ of lists from *Ex. 27*, and consider a second algebra $(Z, [r_0, r_1])$ over the same functor. A homomorphism $f : \mathsf{List}(A) \to Z$ from the former into the latter must satisfy

$$f(\mathsf{Nil}) = r_0,$$
$$f(\mathsf{Cons}(a, l)) = r_1(a, f(l)).$$

Hence, f is necessarily a fold over a list with recursive components r_0 and r_1. Incidentally, this proves that f is unique, making $(\mathsf{List}(A), [\mathsf{Nil}, \mathsf{Cons}])$ an initial object in the category of algebras over the polynomial endofunctor $1 + A \times X$.

The notion of (co)algebra over a functor can be generalized to the sphere 2-categories, defining the notion of (co)algebra over a 2-endofunctor. The basic concepts stay the same but the commutativity of the diagrams defining (co)algebra homomorphisms is relaxed into lax-commutativity. A square diagram of 1-cells is lax-commutative if there exists a 2-cell that carries the composition of the top and the right edge onto the composition of the left and bottom edge, as in ADDIM. Once again we refer to [GLD$^+$24] for further information.

We conclude the section with the following proposition ([GLD$^+$24]), which shows that actegorical strong endofunctors induce 2-endofunctors over parametric categories. We will use the theorem to study parametric versions of the endofunctors in *Ex. 27* and *Ex. 28*.

Proposition 31. *Suppose* $(\mathcal{C}, \bullet)$ *is an* $\mathcal{M}$-*actegory and* $F : \mathcal{C} \to \mathcal{C}$ *is an actegorical endofunctor with strength* σ. *Then,* F *induces a 2-functor* $\mathbf{Para}(F) : \mathbf{Para}_\bullet(\mathcal{C}) \to \mathbf{Para}_\bullet(\mathcal{C})$.

Proof. Define $\mathbf{Para}(F)$ so that:

1. $\mathbf{Para}(F)$ acts like F on objects $A : \mathcal{C}$;
2. $\mathbf{Para}(F)(f) = P \bullet F(A) \xrightarrow{\sigma_{P,A}} F(P \bullet A) \xrightarrow{F(f)} F(B)$ for all $(P, f) : \mathbf{Para}_\bullet(\mathcal{C})(A, B)$;
3. $\mathbf{Para}(F)$ leaves reparametrizations unchanged.

$\square$

2.1.2   Integral transform

Remark 32. In accordance with the notation of [DV22], we use $[A, B]$ to represent the set of $A \to B$ functions, where A and B are sets.

Suppose $(R, \oplus, \otimes)$ is a commutative semiring. An integral transform is a transformation that carries a function in $[W, R]$ to a function in $[Z, R]$ following a precise chain of steps. Integral transforms[2] have been introduced by [DV22] to provide a single formalism able to describe both dynamic programming and GNNs. Integral transforms can be encoded as polynomial spans.

Definition 33 (Polynomial span). *A polynomial span is a triplet $(i : X \to W, p : X \to Y, o : Y \to Z)$ of morphisms in* **FinSet**[3]. *i is known as input, p as process, o as ouput. W is known as input set, X as argument set, Y as message set, Z as ouput set. We also ask that the fibers of p have total orderings[4]. The polynomial span (i, p, o) can be graphically represented as a diagrams such as the one of ADDIM.*

Definition 34 (Integral transform). *Let $(R, \oplus, \otimes)$ be a commutative semiring. Let $(i : X \to W, p : X \to Y, o : Y \to Z)$ be a polynomial span. The associated integral transform is the triplet $(i^* : [W, R] \to [X, R], p_\otimes : [X, R] \to [Y, R], o_\oplus : [Y, R] \to [Z, R])$, where:*

1. *i^* is the pullback mapping $f \mapsto i \mathbin{\text{\fontencoding{OT1}\selectfont\char59}} f$;*
2. *$p_\otimes$ is the argument pushforward mapping*

$$p_\otimes(a)(u) = \bigotimes_{e \in p^{-1}(u)} a(e);$$

3. *$o_\otimes$ is the message pushforward mapping*

$$o_\oplus(m)(v) = \bigoplus_{e \in o^{-1}(v)} m(e).$$

The integral transform $(i^*, p_\otimes, o_\oplus)$ can be represented by the diagram in ADDIM.

Remark 35. While defining i^* is quite straight-forward, defining $p_\otimes$ and $o_\otimes$ is more difficult because the different direction of the arrows implies that the underlying functions p and o must be inverted before considering the associated pullbacks. However, inverting non-invertible functions yields functions into the powersets of the original domains. Moreover, if we want to preserve the original multiplicity of arguments and messages, we have to construct inverses that go into the sets of multisets over the original domains. Hence why we need $\otimes$ and $\oplus$ to aggregate results over such multisets. The significance these steps will be clarified later on in this chapter.

---

[2]The label integral transform refers to the fact that similar ideas can be used to write a categorical definition for familiar analytical integral transforms ([Wil10]). A similar construct is also used in physics ([EPWJ80]).

[3]**FinSet** is the category of finite sets and functions.

[4]Neither we nor [DV22] use this requirement but, as stated in the original paper, the requirement is useful to support functions with non-commuting arguments.

2.2 Categorical deep learning

The optic-based framework we presented in the last chapter provides a structured general-purpose compositional framework for gradient-based learning. However, such great versatility has a price: optics are unable to guide the architectural design of our models. It has been shown times and times again that a better architecture makes as much of a difference in machine learning as an algorithm with better asymptotic cost does in classical computer science. Therefore, finding a principled mathematical framework able to guide such architectural choices is of paramount importance.

Designing an architecture is almost synonymous to imposing a set of constraints: for instance, convolutional layers are notoriously equivalent to translationally equivariant linear layers (ADDREF). Hence, it makes sense to inform the choice of architecture by informing the choice of constraint. Among the various attempts at developing a principled theory of machine learning constraints, geometric deep learning (see e.g. [BBCV21]), or GDL is particularly relevant. Categorical deep learning, or CDL, which is the protagonist of this section, evolved as categorical generalization of GDL.

2.2.1 From GDL to CDL

Geometric deep learning focuses on equivariance constraints defined with respect to group actions, according to the following definition.

Definition 36 (Group action equivariance and invariance)**.** *Let G be a group and let $(S, \cdot)$ and $(T, *)$ be G-actions. A function $f : S \to T$ is equivariant with respect to the aforementioned actions if $f(g \cdot s) = g * f(s)$ for all $s \in \mathcal{S}$ and for all $g \in \mathcal{G}$. We say that f is invariant if $*$ is the trivial action on T, and thus $f(g \cdot s) = f(s)$ for all s and g.*

The GDL approach prescribes that a desired constraint should be expressed in terms of equivariance with respect to group action so that associated equations can be derived. Solving such equations usually implies tying weights, which reduces the net number of parameters in the models and provides the various advantages (ADDREF, ADDREF, ADDREF) afforded by equivariance.

GLD and has seen a lot of success in contexts where data transformations can be expressed in group-theoretical terms. Regrettably, this is not always the case, as many transformations we can subject the data to are either not invertible or even not compositional ([GLD+24]). Moreover, as highlighted in [GLD+24], although GDL is effective at describing the constraints that a model should implement, it is not always clear how such constraints can be actually be implemented since solving the equations associated with a constraint might be extremely challenging. Categorical deep learning, introduced by [GLD+24], attempts to solve the aforementioned problems using category theory[5] to generalize the notion of equivariance and to tie specific constraints to specific architectures.

Remark 37. At the moment, to the best of our knowledge, [GLD+24] is the only publicly available paper that discusses the ideas of CDL.

---

[5]As stated in [BBCV21], GDL is inspired by the *Erlangen Programme*, which unified geometry around the notion of invariant at the end of the nineteenth century. Since category theory can be seen as an extension of the *Programme* (as we already remarked in the introduction to this thesis), it is only natural to attempt to generalize GDL by categorical means.

The main insight of categorical deep learning is that group actions can be represented as algebras over the group action monads, and that maps that are equivariant with respect to these actions are homomorphisms between these algebras. Hence, GDL can be generalized by taking into consideration (co)algebras over other monads and also generic endofcuntors. Hence, CDL generalizes GDL into a theory of (co)algebras over endofunctors and monads, yielding, in the words of the authors of [GLD+24], as a "theory of all architectures". The last statement might be a bit too optimistic, since the only architecture examined at the moment are recurrent and recursive neural networks, but the results shown in [GLD+24] are very promising and, hopefully, future work will widen the scope of CDL to include derivations for other architectures.

The following proposition and the associated example show how exactly CDL subsumes GDL.

Proposition 38. *Let $(G, e, \cdot)$ be a group. Consider the endomorphism $G \times - : \mathbf{Set} \to \mathbf{Set}$ which maps $S \mapsto G \times X$ and $f \mapsto G \times f$. G can be given a monad structure using the natural transformations η, with components $\eta_S : s \mapsto (e, s)$, and μ, with components $(g, h, s) \mapsto (g \times h, s)$. The monad $(G \times -, \eta, \mu)$ can serve as a signature for G-actions. The actions themselves can be recovered by considering algebras $(S, *)$ for the monad, and, given two actions $(S, *)$ and $(T, \star)$, an associated equivariant map $f : S \to T$ is a $(S, *) \to (T, \star)$ monad algebra homomorphism.*

Proof. It suffices to compare the equations that define group actions and group action invariance with the commutative diagrams in ADDIM and ADDIM. □

Example 39 (Linear equivariant layer)**.** Consider a carrier set $S = \mathbb{R}^{\mathbb{Z}_2}$, which can be seen as a pair of pixels. Consider the translation action $(i * s)(j) = s(i - j)$ of $G = \mathbb{Z}_2$ on S, which can be seen as swapping the pixels. We want to find a linear map $f : S \to S$ which is equivariant with respect to the action. Imposing the equivariance constraints as equations on the entries of the matricial representation $W_f \in \mathbb{R}^{2 \times 2}$ of the map, we can prove that f is equivariant if and only if W_f is symmetric.

2.2.2 (Co)inductive definitions for RNNs

As seen in *Ex. 39*, the formalism of CDL subsumes the formalism of GDL, but the the difference between the two is not a simple matter of notation: CDL is much more general. The most significant piece of novel contribution delineated in [GLD+24] is the use of (co)algebras and (co)algebra homomorphisms over parametric categories to define (co)inductively recurrent and recursive neural networks. (Co)algebras are used to define cells, whereas the associated homomorphisms provide the weight-sharing mechanics used to unroll them. Let us build on *Ex. 27* and *Ex. 30*, as is done in [GLD+24].

Example 40 (Folding recurrent neural network cell)**.** Consider the endofunctor $1 + A \times X : \mathbf{Set} \to \mathbf{Set}$ from *Ex. 27*. Consider the Cartesian action of **Set** on itself and associate the following actegorical strength to the functor: $\sigma_{P,X}(p, \mathrm{inl}) = \mathrm{inl}$ and $\sigma_{P,A}(p, \mathrm{inr}(x, x')) = \mathrm{inr}((p, x), (p, x'))$. Now that the functor is actegorical strong, we can use *Prop. 31* to construct an endofunctor $\mathbf{Para}(1 + A \times X) : \mathbf{Para}_\bullet(\mathbf{Set}) \to \mathbf{Para}_\bullet(\mathbf{Set})$. Consider an algebra $(S, (P, \mathsf{Cell}))$ for this functor. Via the isomorphism $P \times (1 + A \times X) \cong P + P \times A \times X$, we deduce that $\mathsf{Cell} = [\mathsf{Cell}_0, \mathsf{Cell}_1]$, where $\mathsf{Cell}_0 : P \to S$ and $\mathsf{Cell}_0 : P \to S$. We can interpret Cell_0 and Cell_1 as folding recurrent neural network cells: Cell_0 provides

the initial state based on its parameter and Cell_1 takes in the old state, a parameter, and an input, which are then used to return a new state (ADDIM).

Example 41 (Unrolling of a folding recurrent neural network)**.** Use *Prop. 7* to embed the list algebra $(\mathsf{List}(A), [\mathsf{Nil}, \mathsf{Cons}])$ from *Ex. 27* as an algebra over the endofunctor $\mathbf{Para}(1 + A \times X)$ from the previous example. Now consider an algebra homomorphism $(P, f) : (\mathsf{List}(A), [\mathsf{Nil}, \mathsf{Cons}]) \to (S, (P, \mathsf{Cell}))$. Since we are working with algebras over a 2-endofunctor, we also need to specify a 2-cell that makes the diagram in ADDIM lax-commutative. Using the weight-tying reparameterization Δ_P, yields the lax commutative diagram in ADDIM, which uniquely identifies f as the fold function which takes a list of inputs in A and unrolls a folding recurrent neural network that reads such inputs. The weight-tying reparametrization makes sure that each cell of the unrolled network uses the same parameters (see ADDIM for a graphical representation).

The construction in *Ex. 40* and *Ex. 41* constitutes a precise mathematical link between the classical data structure of lists and the machine learning construct of folding RNN. Similarly, [GLD$^+$24] constructs recursive neural networks building upon classical binary trees and, even more interestingly, complete RNNs are build upon the coalgebra of *Ex. 28*, which reveals an unexpected link between RNNs and Mealy machines. This begs the question: if Mealy machines generalize to recurrent neural networks, what do Moore machines generalize to? It is argued in the paper that they generalize to a variant of RNN where different cells (which share the same weight) are used for state update and output production. ADDIM shows how different classical concepts generalize to difference NN cells. Hopefully, more work in this direction will lead to new NN arhictecutres inspired from other classical concepts.

Remark 42 (CDL and optic-based learning)**.** In all the cases discussed above, the (co)algebra homomorphism map returns parametric maps (P, model), which we can interpreted as untrained NN models. We can feed these maps into the $\mathbf{R}_{\mathcal{C}}$ functor associated with a generalized Cartesian reverse differential category[6] to augment them with their reverse derivative. The framework of parametric lenses described in *Sec. 1.2* can then be used to train these networks. CDL and optic-based learning are thus perfectly compatible, and even complement each other.

2.3   Algorithmic alignment: GNNs and dynamic programming

One of them main tenets of neural algorithmic reasoning is *algorithmic alignment* ([XLZ$^+$19]), i.e. i.e. the presence of structural similarities between the subroutines of a particular algorithm and the architecture of the neural network selected to learn such algorithm. Since [XLZ$^+$19] has shown that dynamic programming algorithms align very well with GNNs, and since dynamic programming encompasses a wide variety of techniques algorithms to many different domains, this puts GNNs at the forefront of neural algorithmic reasoning research ([DV22]). However, the exact link between GNNs and dynamic programming has yet to be fully formalized. In this section we present an attempt at such formalization described in [DV22]. Fainlly, we also describe the work in [DvGPV24], which uses

---

[6]The examples illustrated in this section have been developed in **Set**, but we see no reason why they couldn't be specialized to an appropriate CRDC.

the formalism of cocycles to derive better implementations of GNNs that are invariant under various form of asynchrony[7].

2.3.1　Integral transforms for GNNs and dynamic programming

The main link between dynamic programming and GNNs is that dynamic programming itself induces a graph structure: given a problem x, we can break it up into subproblems y of x. Then, we can break up the y's and so on, until we get to base cases. We can thus consider the graph with nodes corresponding to subproblems and edges (y, x) corresponding to the relationships 'y is a subproblem of x '. Finally, we can solve the problem by message passing: the simpler cases are solved first and their solutions are passed as messages along the edges so that they can be used to solve more complex cases. The architecture of a graph neural network usually implements such a message passing dynamic. However, rigorously formulating an isomorphism between the architecture of GNN and the structure an associated dynamic programming algorithm is not easy, the main obstacle being the difference in data type handled by the two mathematical processes: dynamic programming usually deals with tropical objects such as the semiring $(\mathbb{N} \cup \{\infty\}, \min, +)$, while GNNs usually deal with linear algebra over $\mathbb{R}$ ([DV22]).

[DV22] proposes the formalism of integral transforms as the common structure behind both GNNs and dynamic programming. While a full formal proof is not given, the idea is illustrated by showing that both the Bellman-Ford algorithm and the message passing neural network (MPNN) architecture can be expressed with the help of integral transforms. In the formalism of [DV22], the difference in data type is overcome by using the weakest common hypothesis: that the data and associated operations form a semiring.

Bellman-Ford algorithm

The Bellman-Ford (BF) algorithm is one of the most popular dynamic programming algorithms and is used to find the shortest paths between a single starting node and every other node in a weighted graph $G = (V, E)$. We can see every node of the graph as a subproblem and we can see the associated edges as subproblem relationships. Hence, the BF algorithm perfectly falls into place as a candidate for a GNN implementation. The algorithm operates within the tropical min-plus semiring $(R = \mathbb{N} \cup \{\infty\}, \oplus = \min, \otimes = +)$, and the data can be provided as a tuple (d, b, w) of three functions into R. Here, $d : V \to R$ stores the current best distances of the nodes, $b : V \to R$ stores the weights of the nodes, and $w : E \to R$ stores the weights of the edges. d is initalized as the function that maps the initial node to 0 and every other node to ∞. The values of the functions are updated at each step of the algorithm according to the following formula, where $\mathcal{N}_u = \{v \text{ s. t. } (v, u) \in E\}$ represent the one-hop neighborhood of a node u:

$$d_u \longleftarrow \min \left(d_u, \min_{v \in \mathcal{N}_u}(d_v) + w_{v,u} \right).$$

---

[7]While much of the work described in [DvGPV24] does not fall under the umbrella of applied category theory, we still mention it because on its close link with the work of [DV22] and the ideas of algorithmic alignment. Hopefully, future work will explore the intersection between this work and category theory.

[DV22] propose the integral transform encoded by the polynomial span in ADDIM as the supporting structure of the BF algorithm. The functions i, p, and o are defined as follows:

1. $i : (V + E) + (V + E) \rightarrow V + (V + E)$ acts as the identity on the first V and on the second $V + E$ pair, whereas maps the edges of the first E to their sources;
2. $p : (V + E) + (V + E) \rightarrow V + E$ just collapses the two copies of $V + E$;
3. o acts as the target function on the E and as the identity on V.

It is argued in the paper that the whole integral transform acts as step of the algorithm, carrying the data in (d, b, w) to the update function d. Let's examine each step: the input pullback extracts the distances the sources of every edge; the argument pushforward computes the weights of the paths from each node to the neighboring nodes (the weight of each node is treated as the weight of a self-edge in this case); finally, the message pushforward selects the shortest paths among the ones studied by the argument pushforward. Hence, the simple polynomial span in ADDIM successfully encode the whole BP algorithm without any information loss or ad hoc choice.

Message passing neural network

A message passing neural network (MPNN) is a graph neural network described by the following equations ([GSR$^+$17]):

$$m_v^{t+1} = \sum_{w \in \mathcal{N}_v} M_t(h_v^t, h_w^t, e_{v,w}),$$
$$h_v^{t+1} = U_t(h_v^t, m_v^{t+1}),$$

where t represents the time step, and M_t and U_t are learned differentiable functions.

[DV22] argues that an MPNN can be implemented as the integral transform associated with the polynomial span of ADDIM, with an extra MLP, as show in ADDIM. Here,

1. i sends the first E to 1, it acts as source and target on the second E and third E respectively, and it acts as the identity on the fourth E;
2. p collapses four E's into one;
3. o acts as the target function.

In the associated integral transform, i^* gathers graph, features, node features and edge features, $p_\otimes$ "projects" the such features on the edges, the MLP combines them, and, finally, $o_\oplus$ sends them to the right target.

Although not a perfect representation of the target architecture (remember the extra MLP), the polynomial span in ADDIM can be used to inform the design of new architectures which are obtained by simple manipulations of the arrows or objects in the diagram. For instance, [DV22] uses the integral transform formalism they introduce to investigate possible performance improvements on CLRS benchmark tasks ([VBB$^+$22]) brought by considering messages that reduce over intermediate nodes. It is shown that this architecture leads to better average performance on these tasks, which is a result of better algorithmic alignment.

2.3.2 Asynchronous algorithmic alignment

[DvGPV24] derives conditions under which both algorithms and synchronous GNNs are invariant under asynchronous execution. This is relevant because, as stated in the paper, in many dynamic programming tasks modeled by graphs, only small parts of the aforementioned graphs are changed at each step. A synchronously executed GNN that is trained on these tasks must learn the identity function many times over, which leads to brittleness and wasted computational resources. However, many algorithms are invariant under asynchronous execution of some of its steps. Implementing these algorithms as synchronous GNNs that are invariant under various forms of asynchronous execution leads to better algorithmic alignment and thus more efficient learning.

The authors of [DvGPV24] revise the model explored in [DV22] to include a message function ψ that generates messages based on gathered arguments: see ADDIM for the update diagram. Moreover, the authors argue that it is best to consider GNNs where every graph component that has a persistent state is elevated to the status of node, while transient computations should be carried out along edges. The resulting GNN can be described by the diagram in ADDIM, where ϕ is the update function that updates the persistent state of each node, while δ is the function that computes the arguments of the message function.

It is argued in [DvGPV24] that invariance under asynchrony can be modeled by giving both arguments and messages monoidal structures. For instance, let $(M, 1, \cdot)$ be the message monoid and let $(A, 0, +)$ be the argument monoid. Then, if S is the set of persistent states, we can model state update and argument generation as a function $M \times S \to S \times A$ which maps $(m, s) \mapsto (m \bullet s, \delta_m(s))$. One way toget invariance under asynchorny in message aggregation is to define $\bullet$ as a monoidal action of M on S. However, we must make sure that argument generation is compatible with the unitality and associativity equations of the action. This can only happen if $\delta : m \mapsto \delta_m$ is a 1-cocycle.

Definition 43 (1-cocycle)**.** *A map* $\delta : M \to [S, A]$ *is a 1-cocycle if and only if the following are satisfied:*

1. $\delta_1(s) = 0$ for all $s \in S$;
2. $\delta_{n \cdot m}(s) = \delta_n(m \cdot s) + \delta_m(s)$ for all $s \in S$.

Proposition 44. *The state update function δ described above is asynchronous with respect to message passing if and only if it is a 1-cocycle.*

[DvGPV24] also proves the following.

Proposition 45. *Under the hypotheses described above, a single-input message function ψ supports asynchronous invocation if and only if ψ is a homomorphism of monoids.*

We will not describe the whole formalism of [DvGPV24], but we will show (without proof) its implications on GNN archietcture design.

Example 46. Consider the GNN architecture:

$$x'_u = \phi \left(x_u, \bigoplus_{v \in \mathcal{N}_u} \psi(x_u, x_v) \right),$$

where $\oplus$ is a message aggregator function. The authors of [DvGPV24] derive conditions under which this architecture is invariant under asynchronies in message aggregation, node update, and argument generation.

The GNN is trivially invariant under asynchronous message aggregation if messages $(M, 0, \oplus)$ are given a commutative monoidal structure. Invariance under asynchronies in node updates is obtained by selecting an update function ϕ which sastifies the associative law $\phi(s, m \oplus n) = \phi(\phi(s, m), n)$ for all $m, n \in M$ and for all $s \in S$. Finally, invariance under argument generation is obtained if ϕ satisfies the 1-cocycle equations (*Def.* **??**). [DvGPV24] gives sufficient conditions that ensure that the 1-cocycle equations as well as the associative law are true: it suffices that $M = A = S$ and that $\phi = \oplus$ is idempotent.

Chapter 3

Learning Funtcors

3.1 Functor learning

In the parametric lens framework presented above, learning happens at the morphism level. The framework is very convenient because its categorical structure hides away the automatic differentiation machinery needed for backpropagation, and thus we can just compose the parametric lenses as modules, so that we can form any network we wish. However, parametric lenses cannot be used to abstract away computations and extract the schema of the model, as they inherently carry out computations. Moreover, the generality of the parametric lens framework means that no tools are developed to study and exploit the inherent structure of the data. Finally, parametric lenses do not offer high level tools to deal with invariance ad equivariance constraints.

It has been suggested by [Gav19] and [SY21] that the key to link different layers of abstraction and to exploit the structure of the data is to learn functors instead of morphisms. [Gav19] and [VSP$^+$17] also use functor learning to provide tools that implement abstract constraints on neural networks.

3.1.1 Functors to Separate Layers of Abstraction

The author of [Gav19] takes inspiration from categorical data migration - a field which is located at the intersection of category theory and database theory - to create a categorical framework for deep learning that separates the development of a machine learning process into a number of key steps. The different steps operate on different levels of abstraction and are linked by functors.

Schemas

The first step in the learning pipeline proposed by [Gav19] is to write down the bare-bones structure of the model in question. One can do that using a directed multigraph G, where nodes represent data and edges represent neural networks interacting with such data. Constraints can be added at this level in the form of a set $\mathcal{X}$ of equations that identify parallel paths (see e.g. ADDIM, whose multigraph represent a cycleGAN model).

We can now consider the most abstract categorical representation of such model: its schema.

Definition 47 (Model schema)**.** *The schema of a model represented by a multigraph G is the freely generated category* **Free**(G).

Such schema does not contain any data nor does it do any computation, but it encodes the bare-bones structure of the model. We can take the module **Free**$(G)/\sim$, where $\sim$ is the path congruence relation induced by the equations in $\mathcal{X}$. Depending on the context, the word schema will also refer to such constrained category.

Architectures and models

Given the schema **Free**(G), we can choose an architecture for the model. What we mean by choosing architecture is assigning to each node a Euclidean space and to each morphism a parametric map, which represents an untrained neural network.

Definition 48 (Model architecture)**.** *Let* Free(G) *be a model schema. An architecture for such schema is a functor* Arch : Free$(G) \to$ **Para**(**Smooth**).

The Euclidean spaces Arch maps objects to might be intuitively interpreted as the spaces the data will live in, but it is best to consider data outside the **Para** machinery and in the simpler **Set** category, as this allows for better compartmentalisation. Thus, [Gav19] also defines an embedding functor, which agrees with Arch on objects but exists independently of it.

Definition 49 (Model embedding)**.** *Let* Free(G) *be a model schema and let* Arch *be a chosen architecture. An embedding for such schema is a functor* $E : |$Free$(G)| \to$ **Set** *which agrees with* Arch *on objects*[1].

Consider the function $\mathfrak{p} : (P, f) \mapsto P$, which takes the parameter space out of a parametric map in **Para**(**Smooth**). We can use it to define a function

$$\mathfrak{P} : \mathsf{Arch} \mapsto \prod_{f : \mathsf{Gen}_{\mathbf{Free}(G)}} \mathfrak{p}(\mathsf{Arch}(f)),$$

where $\mathsf{Gen}_{\mathbf{Free}(G)}$ is the set of generating morphisms of the free category on the multigraph G. The function $\mathfrak{P}$ takes and architecture and returns the parameter space. Given $\mathfrak{P}$, we can define parameter specification function.

Definition 50 (Parameter specification function)**.** *Let* Free(G) *be a model schema and let* Arch *be a chosen architecture. A parameter specification function is a function* PSpec *which maps a pair* (Arch, p) *- comprised of an architecture* Arch *and some* $p : \mathfrak{P}($Arch$)$ *- to a functor* Model$_p$: **Free**$(G) \to$ **Smooth***. The functor* Model$_p$ *takes the model schema and returns its implementation according to* Arch*, partially applying* p_f *to each* Arch(f)*, so that we obtain an actual smooth map.*

The functor Model$_p$ can be seen as a bridge between the parametric notion of untrained neural network and the notion of neural network as a smooth map, which makes most sense after training. To better understand the relationship between Arch and Model$_p$ see ADDIM.

---

[1]The reason why the domain E is the discretised schema $|$Free$(G)|$ instead of the original schema Free(G) will be clear once we define datasets.

Datasets and concepts

Now, if we hope to train the model we have defined, we will need a dataset. [Gav19] suggests that a dataset should be represented as a subfunctor of the embedding functor:

Definition 51 (Dataset). *Let E be a model embedding. Then, a dataset is a subfunctor D_E : $|\mathbf{Free}(G)| \to \mathbf{Set}$ which maps every object A of the discretised free category to a finite subset $D_E(A) \subseteq E(A)$.*

Remark 52. The reason why the author of [Gav19] chooses to define E and D_E on discretized categories is because it often happens in practical machine learning that the available data is not paired. In these cases, it would be meaningless to provide an action on morphisms because they would end up being incomplete maps.

Given a node A of G, we have associated to A a Euclidean space $E(A)$, e.g. $\mathbb{R}^n$, and a dataset $D_E(A)$. It makes sense to define another set $\mathfrak{C}(A)$ such that $D_E(A) \subseteq \mathfrak{C}(A) \subseteq E(A)$. A dataset may be considered a collection of instances of something more specific than just vectors; for instance, if we have a finite dataset of pictures of horses, we are clearly interested in the concept of horse, i.e. in the set of all possible pictures of horses, which is much larger than our dataset but still much smaller than the vector space used to host such pictures. The set $\mathfrak{C}(A)$ is the concept represented by $D_E(A)$: in the aforementioned example, $\mathfrak{C}(A)$ might be the set of all images representing horses. Hence, we can call $\mathcal{C}$ the concept functor.

Definition 53 (Concept functor). *Given a schema $\mathbf{Free}(G)/\sim$, an embedding E and a dataset D_E, a concept associated with this information is a functor $\mathfrak{C} : \mathbf{Free}(G)/\sim \to \mathbf{Set}$ such that, if $I : |\mathbf{Free}(G)| \to \mathbf{Free}(G)/\sim$ is the inclusion functor, $D_E \subseteq I; \mathfrak{C} \subseteq E$.*

As [Gav19] states, $\mathfrak{C}$ is an idealization, but it is a useful idealization as it represent the goal of the optimization process: given a dataset $D_E : |\mathbf{Free}(G)| \to \mathbf{Set}$, we wish to learn the concept functor $\mathfrak{C} : \mathbf{Free}(G)/\sim \to \mathbf{Set}$. Total achievement of such goal is clearly impossible as, even in the simplest cases such as linear regression on linearly generated data, finite arithmetics and the finite nature of the learning iteration prevent us from obtaining a perfect copy of the generating function. Nevertheless, we will hopefully design an optimization process which makes the learning iteration converge towards such ideal goal.

Optimization

Now that we know what the optimization goal is, we can define what a task is. The task formalism brings together what has been defined in this section in an integrated fashion.

Definition 54 (Task). *Let G be a directed multigraph, let $\sim$ be a congruence relation on $\mathbf{Free}(G)$ and let $D_E : |\mathbf{Free}(G)| \to \mathbf{Set}$ be a dataset. Then, we call the triple $(G, \sim, D_E)$ a task.*

Once we are assigned a machine learning task $(G, \sim, D_E)$, we have to choose an architecture, an embedding and a concept compatible with the given multigraph, equations and dataset. Then, we specify a random initial parameter with an appropriate parameter specification function. Now we can choose and optimizer, but we must be careful to choose an appropriate loss function. The loss function should incorporate both an architecture specific loss and a path equivalence loss. The former penalizes wrong predictions while the latter penalizes violations of the constraints embodied by $\sim$.

Definition 55 (Path equivalence loss)**.** *Let* $(G, \sim, D_E)$ *be a task. Let* Model_p *be an associated model. Then, if* $f \sim g : A \to B$ *in* G, *we define the path equivalence loss associated with* f, g *and* Model_p *as*

$$\mathcal{L}_{\sim}^{f,g} = \mathbb{E}_{a \sim D_E(A)}[\|\mathsf{Model}_p(f)(a) - \mathsf{Model}_p(g)(a)\|].$$

Definition 56 (Total loss)**.** *Let* $(G, \sim, D_E)$ *be a task. Let* Arch *be an associated architecture, let* Model_p *be an associated model, and let* $\mathcal{L}'$ *be an architecture specific loss. Suppose* γ *is a non-negative hyperparameter. Then, we define the total loss associated with the task, the architecture, the model, and the hyperparameter as*

$$\mathcal{L} = \mathcal{L}' + \sum_{f \sim g} \mathcal{L}_{\sim}^{f,g} \tag{3.1}$$

We can now proceed as usual, computing the loss on the dataset for a number of epochs and updating the parameter p each time. Notice that $\mathcal{L}$ implicitly depends on p because each $\mathcal{L}_{\sim}^{f,g}$. Thus, the explicit formula for the loss changes each time the parameter is updated.

Product Task

It is important to notice that, while the learning iteration employed by [Gav19] is nothing new, the functor approach is actually novel, in that the usual optimization process is used to explore a functor space instead of a simple morphism space. The main advantage offered by this procedure is that different layers of abstraction are separated, which allows greater expressive power when defining tasks and solving them.

For instance, it is shown in [Gav19] that changing the dataset functor can result in semantically different networks and tasks even if we keep the same schema. The example shown in the paper is the following: if we take the cycleGAN schema (see ADDIM) and we pair it with a cycleGAN dataset, we obtain a cycleGAN network. Here, by cycleGAN dataset, we mean a dataset where A and B contain data which is essentially isomorphic, such as pictures of horses and zebras. The semantics of the task thus the following: *learn maps that turns horses into zebras and vice versa.* However, if we select a dataset where A consists in pictures which contain two elements X and Y, and B contains separate images of X and Y, then the semantics of the task become *learn how to separate X from Y.* For example, [Gav19] shows how to use the CelebA dataset to train a neural network able to remove glasses from pictures of faces, or even insert them (see image ADDIM).

This example is especially relevant to the present discussion because it shows how relevant the categorical structure of **Set** can be to machine learning problems. We can interpret pairs *(face, glasses)* as elements of the Cartesian product of the set of faces and the set of glasses. The set of pictures of faces with glasses can instead be considered another categorical product of the aforementioned sets. Then, one interpretation of the task is the following: *find the canonical isomorphisms between two categorical products.* This task, first introduced by [Gav19], is known as product task.

3.1.2 Categorical Representation Learning

While the functorial approach described above provides a categorical interpretation of datasets, namely as functors into **Set**, no categorical structure is given to the data itself, besides the trivial notion that the data lives in sets. It is argued in [SY21] that sometimes data can be given a categorical

structure of its own, and preserving such categorical structure makes learning more efficient. This can be done in two steps: (i) functorially embed the data into appropriately defined embedding categories, (ii) learn functors between such embeddings.

Categorical Embeddings

We will illustrate this procedure with the same example employed in the original paper [SY21]: unsupervised translation of the names of chemical elements from English to Chinese. Suppose we have two datasets containing thousands of chemical formulas of various inorganic compounds; the first dataset labels the elements with English labels, while the second dataset labels them in Chinese. These datasets can be given a categorical structure with elements and functional groups being the objects, and bonds between them being the morphisms. Let $\mathcal{C}$ and $\mathcal{D}$ be the resulting categories.

Any category $\mathcal{C}$ can be functorially embedded in the Euclidean space $\mathbb{R}^n$ by considering the vector space category $\mathcal{R}$ associated to the aforementioned vector space (*Def.* **??**). Given a categorical structure to the embedding codomain, it suffices to define a $\mathcal{C} \to \mathcal{R}$ functor which maps $a \mapsto v_a$ and $f \mapsto M_f$. We will use train a neural network to carry out such mapping. This happens on two separate embedding layers: one maps words to vectors, while the other maps relations to matrices.

Remark 57. It is worth noticing that a single matrix M can be in many hom-sets of $\mathcal{R}$, and thus we can have a number of different pairs of vectors linked by the same morphism.

Concurrence Statistics

The embedding layer described above needs to be trained on the data to be effectively functorial. The authors of [SY21] suggest using concurrence statistics and negative sampling to make sure that the embedded morphisms actually represent the same relations between objects elements as the original morphisms between the original objects. It is in fact posited that concurrence encodes such relations. In the authors' words, 'concurrence does not happen for no reason'.

The training strategy used in the paper is the following: given two embedded words a and b, model the probability of concurrence as $P(a \to b) = \mathsf{sigmoid}(z(a \to b))$, where the logit $z(a \to b)$ is defined as

$$z(a \to b) = F\left(\bigoplus_f v_a^T M_f v_b\right).$$

Here, F is non-linear and $\bigoplus_f$ represents concatenation over all morphisms in $\mathcal{C}$. The idea behind the formula is that it is not possible, given the finite precision of a computer, to compute a matrix M_f such that $v_b = M_f v_a$ exactly. However, we can find M_f such that v_b and $M_f v_a$ are closely aligned. The alignment can be computed as $v_a^T M_f v_b$. The non-linearity F serves as an aggregator for such measurements.

Now, the concurrence probability $p(a, b)$ of two objects $a, b : \mathcal{C}$ can be computed directly from the dataset. Given a negative sampling distribution p_N on objects unrelated to a, we can implement the negative sampling loss

$$\mathcal{L} = \mathbb{E}_{(a,b) \sim p(a,b)} \left(\log P(a \to b) + \mathbb{E}_{b' \sim p(b')}\right),$$

as described in [SY21]. The embedding network can then be trained by maximizing such loss.

Training Functors as Models

Apply the procedure described above to both $\mathcal{C}$ and $\mathcal{D}$ to get categorical embeddings of the datasets. Now, consider a functor $\mathcal{F} : \mathcal{C} \to \mathcal{D}$ that translates between English labels and Chinese labels. Such functor must equate chemical bonds of the same kind, e.g. if f is a covalent bond so is $\mathcal{F}(f)$. It is posited in [SY21] that the action of $\mathcal{F}$ on morphism is sufficient to deduce the action of $\mathcal{F}$ on objects.

The function $\mathcal{F}$ can be precomposed with the $\mathcal{C} \to \mathcal{R}$ embedding and postcomposed with the inverse of the $\mathcal{D} \to \mathcal{R}$ to become a $\mathcal{R} \to \mathcal{R}$ functor. The authors of [SY21] argue that such functor can be represented by a matrix $V_{\mathcal{F}}$ so that $v_{\mathcal{F}(a)} = V_{\mathcal{F}} v_a$ and $M_{\mathcal{F}(f)} = V_{\mathcal{F}} M_f$. Such representation is only meaningful if (i) $V_{\mathcal{F}} M_f = M_{\mathcal{F}} V_{\mathcal{F}}$ for all f, (ii) $V_{\mathcal{F}} M_{\mathrm{id}_a} = M_{\mathrm{id}_{\mathcal{F}(a)}}$ for all a, and (iii) $V_{\mathcal{F}} M_{f \circ g} = V_{\mathcal{F}} M_f V_{\mathcal{F}} M_g$ for all f, g. This is not true for all choices of $V_{\mathcal{F}}$ but, if we choose every v_a to be a unit vector[2], and if we constrain $V_{\mathcal{F}}$ to be orthogonal, (ii) and (iii) are trivially satisfied. The focus can thus be shifted on requirement (i).

Requirement (i) in the previous paragraph can be learned through the following structure loss:

$$\mathcal{L}_{\mathrm{struc}} = \sum_f \| V_{\mathcal{F}} M_f - M_{\mathcal{F}} V_{\mathcal{F}} \|^2 .$$

As the authors remark, this loss is universal, in the sense that it does not depend on any specific object, but acts on the morphisms themselves. While this approach is very elegant and does indeed return a functor, this might not be the functor we expect because $V_{\mathcal{F}}$ is not unique if the M_f happen to be singular. Thus, it is better to integrate the structure loss with a second alignment loss which introduces some supervision to the unsupervised translation task. For instance, if the value of $\mathcal{F}(a)$ is known for a set A for objects, we can define

$$\mathcal{L}_{\mathrm{align}} = \sum_{a:A} \| V_{\mathcal{F}} v_a - v_{\mathcal{F}(a)} \| .$$

Then, the total loss can be written as a weighted sum $\mathcal{L} = \mathcal{L}_{\mathrm{align}} + \lambda \mathcal{L}_{\mathrm{align}}$ of structure loss and alignment loss, where λ is a hyperparameter analogous of the γ that appears in *Eq. 3.1*.

Comparison with Traditional Models

It is shown in [SY21] that the kind of functor learning illustrated in this section can lead to remarkable improvements in efficiency when compared with more traditional sequence to sequence models. In particular, the authors compared a GRU cell model of similar performance as the functorial model described in the paper, noting that the former needed 17 times more parameters than the latter.

The authors also compare their approach with the multi-head attention approach first introduced by [VSP$^+$17] (the title of [SY21] is clearly inspired on the title of [VSP$^+$17]). It is argued in the article that the categorical approach is an improvement over multi-head attention as the matrices M_f are essentially equivalent to the products $Q_f^T K_f$, where Q_f is the query matrix associated to f and K_f is the key matrix associated to f. Categorical representation learning differs from multi-head attention as it does not separate the M_f's into their components, which is useful to learn the matrix $V_{\mathcal{F}}$ and allows us to benefit from functoriality.

---

[2]The paper does not specify what strategy was used. See ADDREF for a possible example of how this might be carried out. Check that this is correct!

3.1.3 Invariance and Equivariance with Functors

Inspired in part by the formalization in [Gav19], [CLLS24] presents a categorical framework that describes neural networks as functors and uses their functorial nature to impose invariance and equivariance constraints. In particular, it is shown that such constraints are helpful in accounting for shift and imbalance in covariates when pooling medical image datasets.

Categorical Structure of Data

The first challenge in using functors to enforce invariance and equivariance constraints is giving data a categorical structure. The strategy employed by [CLLS24] consists in defining a data category $\mathcal{S}$ whose objects s are data points and whose morphisms $f : s_1 \to s_2$ represent differences in covariates.

An example considered in the paper is the following: suppose the objects s are comprised of brain scans and associated information concerning patient age and other covariates. The goal is to develop a model trained to diagnose Alzheimer's disease from the scans. An example of morphism in such data category is $f_x : s_1 \to s_2$, which indicates a difference of x years in age: $s_2.\text{age} = s_1.\text{age} + x$.

Since we are dealing with a classification task, the dataset here has labels. It is important not to include the labels in the data category, else any classifier model would just read such labels instead of learning to predict them. Use the notation $\mathbf{y}_s$ to represent the label associated to s. Since the dataset has labels while the data category must not have them, we feel justified in using the phrase data category in place of the phrase dataset category employed above.

Invariance and Equivariance

Now, if we consider another category $\mathcal{T}$, we can use a functor $F : \mathcal{S} \to \mathcal{T}$ to project $\mathcal{S}$ onto $\mathcal{T}$. Learning such functor instead of a simple map between objects is advantageous because the functoriality axioms imply that F automatically satisfies equivariance constraints: if $g : s_1 \to s_2$,

$$F(g) : F(s_1) \to F(s_2). \tag{3.2}$$

Eq. 3.2 is a categorical generalization of the more usual group theoretical notion of invariance, defined as

$$f(g \cdot s) = g \cdot f(s), \tag{3.3}$$

where $s : S$, $g : G$, G is a group, and S is a G-set. To be precise, *Eq. 3.3* is equivalent to *Eq. 3.2* if $\mathcal{S}$, $\mathcal{T}$ are Borel spaces and g, $F(g)$ are group actions, as highlighted in [CLLS24].

Hence, a functor F is automatically equivariant with respect to any change $g : s_1 \to s_2$ in covariates incarnated as a morphism in the domain category. Invariance with respect to g is not much harder to define: it suffices to impose $F(s_1) = F(s_2)$ and $F(g) = \text{id}_{F(s_1)}$.

Classification Task

A functor F able to satisfy invariance and equivariance can be used as a first step to create a classifier that satisfies such constraints. The architecture proposed by [CLLS24] consists in two modules: (i) an autoencoder whose encoder is $F : \mathcal{S} \to \mathcal{T}$ and whose decoder is $F^{-1} : \mathcal{T} \to \mathcal{S}$; (ii) a functor $C : \mathcal{T} \to \mathbf{Free}(\mathbb{N})$ that actually does the classification. Thus, the whole model admits the

compact representation $F \,\fatsemi\, C : \mathcal{S} \to \mathbf{Free}(\mathbb{N})$ (a diagram representing such structure can be seen in ADDIM).

In the model described above, $\mathcal{T}$ acts as a latent space. The hope is that the latent representation of the data in $\mathcal{S}$ satisfies the given equivariance and invariance constraints, so that the actual classification operated by $\mathcal{C}$ naturally satisfies the requirements as well. This same strategy is used in non-categorical approaches such as ADDREF. The main advantage to the categorical formalization presented by [CLLS24] is that an arbitrary number of covariates can be handled at once without any additional complexity. This is in stark contrast with ADDREF, where at most two covariates can be handled at once, and with ADDREF, where increasing the number of covariates requires a much more complicated training pipeline. The authors argue that the framework can also be easily adapted to regression tasks by replacing $\mathbf{Free}(\mathbb{N})$ with $\mathbf{Free}(\mathbb{R})$.

Remark 58. It only makes sense to consider F^{-1} is F is fully-faithful. In practical applications, this isn't usually a concern.

Training

What is now needed is an algorithm able to learn F and C. [CLLS24] suggests implementing $\mathcal{T}$ as a vector space category (*Def.* **??**) and to train a neural network to embed data points $s : \mathcal{S}$ as vectors $F(s) = v_s$, and covariate morphisms f as matrices w_f. It is often useful to restrict ourselves to representing morphisms using orthogonal matrices, as the latter can be inverted by transposition, which is computationally efficient. As shown in the paper, being able to invert such morphisms offers great benefits.

The embeddings and matrices can now be trained using a linear combination of three separate losses: $\mathcal{L} = \gamma_1 \mathcal{L}_r + \gamma_2 \mathcal{L}_p + \gamma_3 \mathcal{L}_s$, where γ_1, γ_2, and γ_3 are hyperparameters. Here, $\mathcal{L}_r$ is a reconstruction loss, which makes sure that F is invertible and that its inverse accurately reconstructs the original data; $\mathcal{L}_p$ is a prediction loss, which makes sure that $F \,\fatsemi\, G$ accurately predicts the labels of the data; $\mathcal{L}_s$ a structure loss, which makes sure that F acts as a functor and not just a map. In formulae,

$$\mathcal{L}_r = \sum_{s:\mathcal{S}} \|s - (F^{-1} \circ F)(s)\|_2^2,$$

$$\mathcal{L}_p = \sum_{s:\mathcal{S}} \mathsf{crossentropy}(\mathbf{y}_s, (C \circ F)(s)),$$

$$\mathcal{L}_s = \sum_{\substack{s_1, s_2 : \mathcal{S} \\ f : s_1 \to s_2}} \|W_f F(s_1) - F(s_2)\|_2^2.$$

Experimental Results

The authors of [CLLS24] prove the validity of the proposed approach with two interesting experiments: a proof of concept trained on the MNIST dataset, and a working classifier trained on the ADNI brain imaging dataset.

The proposed MNIST model implements equivariance with respect to increments, rotations, and zooming. It is shown in the paper that representing the associated morphisms with orthogonal matrices

allows such morphisms to be inverted and combined in the latent space. A subsequent application of F^{-1} shows the results of the aforementioned manipulation in human-understandable form. Such results are indeed very promising: the authors are able to combine rotations and zooming successfully, even though the network was only trained to apply them separately (see ADDIM) for an example.

The ADNI classifier model also shows very promising results which are on par with state-of-the-art models that do not use categorical tools. The comparison takes place according to accuracy of prediction, maximum mean discrepancy, and adversarial validation.

Bibliography

[BBCV21] Michael M Bronstein, Joan Bruna, Taco Cohen, and Petar Veličković. Geometric deep learning: Grids, groups, graphs, geodesics, and gauges. *arXiv preprint arXiv:2104.13478*, 2021.

[BCS06] Richard F Blute, J Robin B Cockett, and Robert AG Seely. Differential categories. *Mathematical structures in computer science*, 16(6):1049–1083, 2006.

[CCG⁺19] Robin Cockett, Geoffrey Cruttwell, Jonathan Gallagher, Jean-Simon Pacaud Lemay, Benjamin MacAdam, Gordon Plotkin, and Dorette Pronk. Reverse derivative categories. *arXiv preprint arXiv:1910.07065*, 2019.

[CEG⁺24] Bryce Clarke, Derek Elkins, Jeremy Gibbons, Fosco Loregian, Bartosz Milewski, Emily Pillmore, and Mario Román. Profunctor optics, a categorical update. *Compositionality*, 6, 2024.

[CG22] Matteo Capucci and Bruno Gavranović. Actegories for the working amthematician. *arXiv preprint arXiv:2203.16351*, 2022.

[CGG⁺22] Geoffrey SH Cruttwell, Bruno Gavranović, Neil Ghani, Paul Wilson, and Fabio Zanasi. Categorical foundations of gradient-based learning. In *European Symposium on Programming*, pages 1–28. Springer International Publishing Cham, 2022.

[CLLS24] Sotirios Panagiotis Chytas, Vishnu Suresh Lokhande, Peiran Li, and Vikas Singh. Pooling image datasets with multiple covariate shift and imbalance, 2024.

[DV22] Andrew J Dudzik and Petar Veličković. Graph neural networks are dynamic programmers. *Advances in neural information processing systems*, 35:20635–20647, 2022.

[DvGPV24] Andrew Joseph Dudzik, Tamara von Glehn, Razvan Pascanu, and Petar Veličković. Asynchronous algorithmic alignment with cocycles. In *Learning on Graphs Conference*, pages 3–1. PMLR, 2024.

[Ell18] Conal Elliott. The simple essence of automatic differentiation. *Proceedings of the ACM on Programming Languages*, 2(ICFP):1–29, 2018.

[EPWJ80] Michael G Eastwood, Roger Penrose, and Raymond O Wells Jr. Cohomology and massless fields. 1980.

[FJ19]   Brendan Fong and Michael Johnson. Lenses and learners. *arXiv preprint arXiv:1903.03671*, 2019.

[FST19]   Brendan Fong, David Spivak, and Rémy Tuyéras. Backprop as functor: A compositional perspective on supervised learning. In *2019 34th Annual ACM/IEEE Symposium on Logic in Computer Science (LICS)*, pages 1–13. IEEE, 2019.

[Gav19]   Bruno Gavranović. Compositional deep learning. *arXiv preprint arXiv:1907.08292*, 2019.

[Gav24]   Bruno Gavranović. Fundamental components of deep learning: A category-theoretic approach, 2024.

[GLD⁺24]   Bruno Gavranović, Paul Lessard, Andrew Joseph Dudzik, Tamara von Glehn, João Guilherme Madeira Araújo, and Petar Veličković. Position: Categorical deep learning is an algebraic theory of all architectures. In *Forty-first International Conference on Machine Learning*, 2024.

[GSR⁺17]   Justin Gilmer, Samuel S. Schoenholz, Patrick F. Riley, Oriol Vinyals, and George E. Dahl. Neural message passing for quantum chemistry. In Doina Precup and Yee Whye Teh, editors, *Proceedings of the 34th International Conference on Machine Learning*, volume 70 of *Proceedings of Machine Learning Research*, pages 1263–1272. PMLR, 06–11 Aug 2017.

[IKP⁺22]   Borja Ibarz, Vitaly Kurin, George Papamakarios, Kyriacos Nikiforou, Mehdi Bennani, Róbert Csordás, Andrew Joseph Dudzik, Matko Bošnjak, Alex Vitvitskyi, Yulia Rubanova, et al. A generalist neural algorithmic learner. In *Learning on graphs conference*, pages 2–1. PMLR, 2022.

[JR97]   Bart Jacobs and Jan Rutten. A tutorial on (co) algebras and (co) induction. *Bulletin-European Association for Theoretical Computer Science*, 62:222–259, 1997.

[Ril18]   Mitchell Riley. Categories of optics. *arXiv preprint arXiv:1809.00738*, 2018.

[SGW21]   Dan Shiebler, Bruno Gavranović, and Paul Wilson. Category theory in machine learning. *arXiv preprint arXiv:2106.07032*, 2021.

[SY21]   Artan Sheshmani and Yi-Zhuang You. Categorical representation learning: morphism is all you need. *Machine Learning: Science and Technology*, 3(1):015016, 2021.

[VB21]   Petar Veličković and Charles Blundell. Neural algorithmic reasoning. *Patterns*, 2(7), 2021.

[VBB⁺22]   Petar Veličković, Adrià Puigdomènech Badia, David Budden, Razvan Pascanu, Andrea Banino, Misha Dashevskiy, Raia Hadsell, and Charles Blundell. The clrs algorithmic reasoning benchmark. In *International Conference on Machine Learning*, pages 22084–22102. PMLR, 2022.

[VSP+17]   Ashish Vaswani, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan N
           Gomez, Łukasz Kaiser, and Illia Polosukhin. Attention is all you need. *Advances in
           neural information processing systems*, 30, 2017.

[Wil10]   Simon Willerton. Integral transforms and the pull-push perspective, 2010.

[Wis08]   Robert Wisbauer. Algebras versus coalgebras. *Applied Categorical Structures*, 16(1):255–
           295, 2008.

[WZ21]   Paul Wilson and Fabio Zanasi. Reverse derivative ascent: A categorical approach to
           learning boolean circuits. *arXiv preprint arXiv:2101.10488*, 2021.

[WZ22]   Paul Wilson and Fabio Zanasi. Categories of differentiable polynomial circuits for machine
           learning. In *International Conference on Graph Transformation*, pages 77–93. Springer,
           2022.

[XLZ+19]   Keyulu Xu, Jingling Li, Mozhi Zhang, Simon S Du, Ken-ichi Kawarabayashi, and Stefanie
           Jegelka. What can neural networks reason about? *arXiv preprint arXiv:1905.13211*, 2019.

Acknowledgements

I wish to acknowledge the essential role my advisor Professor F. Zanasi has had in guiding me through the process of writing this thesis. He introduced me to applied category theory, to machine learning, and to the world of academic research, and I will be forever grateful for it. I also wish to thank my family and my friends, who supported my throughout this journey, and without whom all of this would not have been possible. To all of you, my most sincere gratitude.